

Stored Is Not Supported: Typed Provenance and Assertion Guardrails for Persistent AI Agents

Jun He
OpenKedge.io

Deying Yu
OpenKedge.io

Abstract

Persistent AI agents construct autobiographical state through reflection, retrieval, and consolidation. Persistence changes availability, not epistemic standing: stored or retrieved material is not thereby supported. Untrusted inputs, prompt injections, and model inferences can therefore enter persistent state and later be presented as agent history or user commitments. We specify typed provenance and assertion guardrails for *autobiographical assertion boundedness*, a system-relative release property requiring governed statements about the agent, user, or named relationships to satisfy accepted-evidence, temporal-validity, and disclosure policies. A typed provenance graph separates origin, dependency lineage, epistemic role, validity, and disclosure scope. A resolver evaluates authorized state projections and returns one evidential status, orthogonal conflict, staleness, and withholding flags, and a protected decision witness. A generate–verify–revise mediator then checks candidate semantic units before release and renders policy-authorized status responses. Under explicit assumptions about extraction, predicate correctness, resolution soundness, view declassification, and channel mediation, we prove a conditional assertion-boundedness contract. In an executable suite of 24 hand-authored conformance cases, typed mediation passed none of 19 unsafe opportunities unqualified while preserving all five supported controls. The flat/prior and source-tag comparison rules released 19/19 and 18/19 unsafe candidates, respectively. These results validate the encoded resolver and mediator obligations; they do not constitute an end-to-end evaluation of language models or retrieval systems.

1 Introduction

Autonomous AI agents increasingly maintain long-term state across extended interactions, employing hierarchical memory architectures that combine periodic reflection, semantic retrieval, and background consolidation [1–4]. Persistence expands what an agent can retrieve without establishing what

it may assert. Benchmarks such as LoCoMo evaluate factual retrieval over extended conversational horizons [5]; retrieval alone cannot confer epistemic entitlement. Storing or retrieving an unverified string or model inference does not authorize its presentation as an authentic part of the agent’s autobiographical history.

Consider an agent that ingests an untrusted document containing an embedded indirect instruction asserting that the user prefers public data disclosure. During background memory consolidation, a language model may distill this unverified instruction into a persistent user profile. In a subsequent session, the agent retrieves the isolated summary and asserts with first-person authority: “You asked me to share this information publicly.” Through recurrent retrieval and re-summarization loops, such ungrounded claims can self-corroborate, while stronger successor models may embellish them with plausible parametric details. Prompt-injection evaluations show that untrusted inputs can steer agent reasoning [6]; in a persistent architecture, a resulting unsupported claim can survive the originating context and influence later agent-specific responses.

Common mitigations—such as tracking source URLs, scoring model confidence, or prompting models for self-reflection—do not jointly enforce the policy dimensions considered here. A cryptographically authenticated source may supply inaccurate information; multiple retrieved records may trace back to a single upstream root, manufacturing an illusion of consensus; historically valid preferences can expire; and legitimately supported facts may be restricted from specific recipients. Crucially, procedural entitlement requires decoupling orthogonal operational dimensions: state admission (reachability from an authoritative head), evidential support (conformance to explicit dependency and derivation policies), agent stance (recorded belief versus attributed report), temporal validity (active versus expired), and disclosure authorization (recipient and purpose access control). None of these checks decides metaphysical truth; rather, they enforce auditable procedural bounds over what an agent is permitted to claim.

We address the central question:

Can a persistent agent distinguish what it stores or its underlying model generally knows from what this specific agent has evidence, authority, and valid scope to claim?

We formalize *autobiographical assertion boundedness*, a system-level release property governing outward statements concerning the agent, its user, or named relationships. Rather than attempting the intractable task of pruning a model’s latent parametric knowledge, our architecture treats storage and support as separate judgments by regulating transitions between explicit epistemic roles. Model priors, raw retrieved inputs, admitted evidence, derived interpretations, and recorded beliefs remain distinct persisted objects. Outward autobiographical assertions are mediated through an explicit release gate: unqualified content assertions must bind to admitted, policy-supported, current, and disclosure-authorized claims at an exact state head, whereas resolution-status responses must safely render internal resolver outcomes without leaking protected metadata.

Separating storage from support requires more than a well-formed memory store: an opaque model can hallucinate biographical claims even when the store is flawless. We therefore establish three distinct enforcement surfaces:

- **State admission:** Governs which claim-bearing objects become authoritative in the agent’s state store;
- **Authorized projection:** Filters accepted state into task- and recipient-specific views prior to model context assembly;
- **Assertion mediation:** Intercepts candidate responses at governed output channels to verify content support and disclosure safety before release.

Together, these surfaces prevent persistence from silently converting availability into assertion authority.

Accepted-state boundary. Transactional state mechanisms govern which versioned objects become reachable at an authoritative head [7, 8]. This paper takes that accepted head as an authenticated input and addresses a different question: which reachable claims support a particular outward statement for a particular recipient and time. Atomic storage therefore remains an assumption at the lower boundary, not a contribution or evaluation target of this work. Reachability supplies candidates for resolution; it does not establish support by itself.

Contributions. This paper makes four contributions:

1. **Typed autobiographical claim model.** We formulate a typed provenance graph over evidence and claim vertices, with policy-relative root recognition and role pro-

motion that separate origin, dependency lineage, epistemic role, validity time, and disclosure scope (Section 2).

2. **Procedural resolution and forward correction.** We design a typed resolver that evaluates authorized projections to return a single evidential status alongside orthogonal conflict, staleness, and withholding flags, generating a minimized, auditable decision witness and supporting dependency-aware invalidation (Section 3).
3. **Assertion mediation pipeline.** We formulate a generate–verify–revise release gate and prove a conditional contract result for content assertions and recipient-safe status renderings under emission-head verification (Section 4).
4. **Conformance validation and evaluation design.** We provide an executable conformance suite covering 24 structured cases across six threat tracks, together with metrics and a separate protocol for future end-to-end evaluation (Section 5).

Scope and boundaries. Our guarantees apply to declared autobiographical claims on mediated output channels over an intact accepted head. The executable artifact tests structured resolver and mediator behavior; it does not implement natural-language semantic extraction, generalized witness access control, generalized view-equivalence checks, full provenance-graph traversal, canonical assertion-artifact hashing, or target-origin admission checks. The mechanism also does not implement lower-level concurrency or crash recovery, certify external truth, audit unmediated out-of-band channels, or make claims regarding machine consciousness or subjective identity.

2 Epistemic State Model and Threat Boundary

We model the epistemic state of an autonomous agent over an abstract, authenticated state-transition interface. The results below are conditional on four environment assumptions:

- **(E1) Authoritative state ordering:** A monotonically versioned branch head h identifies a finite reachable set of accepted, immutable epistemic objects, separating them from uncommitted, speculative, or rejected candidates.
- **(E2) Versioned dependency closure:** Every evidence, claim, and derivation identifier within the accepted state resolves deterministically to an immutable content version or an explicit, authenticated unavailable marker.
- **(E3) Immutable policy binding:** Epistemic admission, projection, and release policies κ are versioned, authenticated, and bound to execution contexts.

- **(E4) Closed channel mediation:** Every governed out-bound communication channel routes candidate agent emissions through the assertion mediator without un-monitored bypass paths.

Append-only logs, Merkle DAGs, and transactional stores can realize E1–E2 (e.g., [7, 8]). Our theoretical formulation treats the accepted head h as an abstract authenticated input. The model therefore focuses on epistemic classification, dependency topology, and warranted assertibility.

2.1 Covered Autobiographical Language

Let $\mathcal{L}_{bio} \subset \mathcal{L}$ denote the covered formal sublanguage of structured autobiographical propositions. A proposition $\phi \in \mathcal{L}_{bio}$ is a typed tuple specifying a subject, predicate, argument terms, temporal interval, and epistemic modality, where:

1. The subject designates the agent itself, its authenticated principal (user), or a named relational entity within their shared operational environment; and
2. The predicate characterizes an experiential interaction, stated preference, granted permission, relational commitment, task capability, or historical event.

Under this formulation, autobiographical standing is semantic and system-relative rather than purely grammatical. Assertions such as “the user requested plain-text formatting” or “the agent holds active consent for calendar inspection” belong to $\mathcal{L}_{bio}$ irrespective of surface syntactic framing or pronoun choice.

Deployments parameterize $\mathcal{L}_{bio}$ via an explicit domain signature $\Sigma = \langle \mathcal{P}_{pred}, \mathcal{A}, \mathcal{M} \rangle$ defining permitted predicates $\mathcal{P}_{pred}$, argument sorts $\mathcal{A}$, and modal operators $\mathcal{M}$. General world knowledge (e.g., mathematical truths or physical laws) remains outside $\mathcal{L}_{bio}$ unless explicitly framed as an agent-specific episodic experience or commitment (e.g., “the agent retrieved this document during yesterday’s session”).

2.2 Model Parametric Capacity vs. Accepted Agent Claims

A persistent agent couples an underlying generative model substrate m with an accepted state head h . We distinguish two information spaces with different evidential status:

1. **Accepted Agent Claim Space:** For any state head h , let

$$\mathcal{C}^{agent}(h) = \text{ClaimsReachable}(h) \quad (1)$$

denote the finite, enumerable set of typed claim objects accepted into the agent’s authoritative state at h .

2. **Latent Model Parametric Space:** Let

$$\mathcal{I}_m^{model} = \{x \mid x \text{ is latent, contextual, or retrieved information accessible to } m\} \quad (2)$$

represent the opaque space of statistical correlations, pre-trained weights, and prompt-context associations accessible to model m . The release mechanism does not assume that this space is available as an enumerable, authenticated set of claims.

Let $\Phi^{agent}(h) = \{\phi(v) \mid v \in \mathcal{C}^{agent}(h)\}$ be the propositions carried by accepted claim objects. Rather than attempting to prune or retrain $\mathcal{I}_m^{model}$, the release contract uses the following non-entitlement principle:

$$\text{CanGenerate}_m(\phi) \not\Rightarrow \phi \in \Phi^{agent}(h). \quad (3)$$

The model’s ability to generate a proposition does not give that proposition standing in the agent’s historical state. Consequently, upgrading or replacing the model substrate ($m_1 \rightarrow m_2$) under a fixed state head h may alter linguistic fluency, general reasoning, or heuristic planning, but leaves the accepted autobiographical claim set $\mathcal{C}^{agent}(h)$ invariant.

2.3 Decoupling Truth, Provenance, and Warranted Assertibility

Semantic truth and procedural support are distinct predicates. Let $\text{True}(\phi)$ denote an external truth predicate that the mechanism does not decide. For a claim vertex v with $\phi(v) \in \mathcal{L}_{bio}$, evaluated at head h under policy κ , the model assumes none of the following implications:

$$\text{Admitted}(v, h) \not\Rightarrow \text{True}(\phi(v)), \quad (4)$$

$$\text{Authenticated}(v) \not\Rightarrow \text{True}(\phi(v)), \quad (5)$$

$$\text{Supported}(v, h; \kappa) \not\Rightarrow \text{True}(\phi(v)), \quad (6)$$

$$\neg \text{Supported}(v, h; \kappa) \not\Rightarrow \neg \text{True}(\phi(v)). \quad (7)$$

Here, $\text{Admitted}(v, h)$ denotes reachability within the state store; $\text{Authenticated}(v)$ confirms cryptographic custody of the recorded origin; and $\text{Supported}(v, h; \kappa)$ denotes satisfaction of the governing evidence and derivation policy.

Table 1 summarizes these decoupled dimensions. Cryptographic signatures guarantee message integrity, not informant veracity; retrieval frequency reflects index topology rather than independent epistemic corroboration; model generation confidence measures probability calibration rather than evidential proof; and admitting a statement into memory records the event of its receipt rather than certifying the objective reality of its content.

2.4 Epistemic State Spaces and Provenance DAGs

We formalize the accepted epistemic state at head h as a labeled, directed acyclic provenance graph:

$$\mathcal{G}_h = (\mathcal{V}_h, \mathcal{E}_h, \lambda), \quad (8)$$

where each vertex $v \in \mathcal{V}_h$ represents an immutable epistemic entity (a raw evidence artifact or a typed claim). A directed

Table 1: Decoupled epistemic dimensions, formal properties, and operational limits.

Epistemic Dimension	Formal Question	Operational Limit
State Admission	Is the entity reachable at authoritative head h ?	Does not establish objective truth or ongoing validity
Source Authentication	Did the declared principal cryptographically author the payload?	Does not establish informant honesty, accuracy, or source independence
Evidential Support	Does the entity’s provenance DAG satisfy release policy κ ?	Does not guarantee empirical infallibility or global completeness
Agent Stance	Is the entity typed as an observation, attributed report, or hypothesis?	Does not certify external correctness
Disclosure Authority	Does recipient p and purpose u satisfy scope predicate $\mathbb{S}$?	Does not imply public domain availability
Semantic Truth	Does the proposition accurately describe external physical reality?	Not decided by this mechanism; depends on external facts and source reliability

edge $(u, v) \in \mathcal{E}_h$ points from a dependency u to an entity v that depends on it.

Let $\mathcal{L}_{bio}^\perp = \mathcal{L}_{bio} \cup \{\perp\}$, where $\perp$ marks an evidence object that carries no autobiographical proposition, and let $\mathcal{P}_{scope} = 2^{\mathcal{P}_{principal} \times \mathcal{U}_{purpose} \times \mathcal{K}_{task}}$ be the authorization-scope domain. The valuation function $\lambda : \mathcal{V}_h \rightarrow \mathcal{T} \times \mathcal{L}_{bio}^\perp \times \mathcal{I}_{time} \times \mathcal{P}_{scope}$ assigns each vertex its structural epistemic properties:

$$\lambda(v) = (\tau(v), \phi(v), \mathbb{T}(v), \mathbb{S}(v)), \quad (9)$$

where:

- $\tau(v) \in \mathcal{T}$ designates the explicit epistemic role;
- $\phi(v) \in \mathcal{L}_{bio}^\perp$ is the structured proposition carried by v , or $\perp$ for non-propositional evidence;
- $\mathbb{T}(v) = [t_{start}, t_{end}] \in \mathcal{I}_{time}$ is its validity interval over the extended time domain $\mathbb{R}$; and
- $\mathbb{S}(v) \in \mathcal{P}_{scope}$ defines the authorization scope across principals, purposes, and tasks.

Let $\mathcal{M}_h$ denote the authenticated unavailable markers returned by E2. These markers are dependency references rather than epistemic vertices: they carry no proposition or role, and a transition record binds each marker to the dependency version it replaces.

Root vertices with indegree zero in $\mathcal{G}_h$ represent foundational source objects (e.g., sensor capture logs, signed user messages, imported external documents). Vertices with non-zero indegree bind an immutable derivation descriptor specifying the transition rule, policy version, and evaluator configuration.

2.5 Source Independence via Provenance Graph Disjointness

A primary failure mode in persistent agent architectures is *false corroboration*: multiple retrieved summaries, search snippets, or paraphrased records derived from a single upstream source are treated as independent consensus. In our

formalization, multi-source corroboration policies must evaluate structural disjointness over the provenance graph rather than counting raw records.

Let $\text{Anc}_h(v) = \{u \in \mathcal{V}_h \mid u \rightsquigarrow v\}$ denote the reflexive, transitive ancestors of v , and let $\mathcal{R}(h) = \{u \in \mathcal{V}_h \mid \text{indeg}(u) = 0\}$ be the source roots. The policy-recognized roots of v are:

$$\text{Roots}_\kappa(v, h) = \{r \in \text{Anc}_h(v) \cap \mathcal{R}(h) \mid \text{RootOK}_\kappa(r, h)\}. \quad (10)$$

Definition 1 (Source Independence). *Two evidence or claim entities $u, v \in \mathcal{V}_h$ are policy-independent under κ , denoted $\text{Indep}_\kappa(u, v; h)$, if and only if their terminal source root sets are nonempty and mutually disjoint. Writing $R_u = \text{Roots}_\kappa(u, h)$ and $R_v = \text{Roots}_\kappa(v, h)$:*

$$\text{Indep}_\kappa(u, v; h) \iff R_u \neq \emptyset \wedge R_v \neq \emptyset \wedge R_u \cap R_v = \emptyset. \quad (11)$$

This graph-theoretic definition is conservative: multiple URLs, distinct semantic embeddings, or syntactic paraphrases that converge onto a shared upstream document share source roots and cannot satisfy independent corroboration requirements.

2.6 Deductive Admission and Role Promotion Calculus

Moving an entity across epistemic roles is a governed state-transition judgment rather than an implicit byproduct of language generation. We use the disjoint role partition $\mathcal{T} = \mathcal{X} \dot{\cup} \mathcal{D} \dot{\cup} \mathcal{U}$, where:

- $\mathcal{X}$ contains `ObservedEvent` and `AttributedReport`, the grounded historical roles anchored in direct sensors or authenticated channels;
- $\mathcal{D}$ contains `DeterministicDerivation`, `MemoryInterpretation`, and `RecordedBelief`, which retain their dependencies;

- $\mathcal{U}$ contains ModelPrior, UnverifiedRetrieval, ProbabilisticInterpretation, and Hypothesis, the synthetic or unverified roles.

The judgment $\mathcal{G}_h \vdash_{\kappa} v_s \xrightarrow[D]{\text{prom}} v_t : r_t$ means that policy κ permits a fresh target v_t with role r_t and exact dependency set $D \subseteq \mathcal{V}_h \cup \mathcal{M}_h$. Admission at the successor head adds v_t , adds an edge (u, v_t) for each $u \in D \cap \mathcal{V}_h$, and retains any marker in $D \cap \mathcal{M}_h$ in the transition record. Two illustrative rules are:

$$\frac{v_s \in D \subseteq \mathcal{V}_h \cup \mathcal{M}_h, \quad \tau(v_s) = \text{AttributedReport} \quad \text{Fresh}_h(v_t), \quad \text{EvalOK}_{\kappa}(D)}{\mathcal{G}_h \vdash_{\kappa} v_s \xrightarrow[D]{\text{prom}} v_t : \text{MemoryInterpretation}} \quad (12)$$

$$\frac{v_s \in D \subseteq \mathcal{V}_h \cup \mathcal{M}_h, \quad \tau(v_s) = \text{ModelPrior} \quad \text{Fresh}_h(v_t)}{\mathcal{G}_h \vdash_{\kappa} v_s \xrightarrow[D]{\text{prom}} v_t : \text{Hypothesis}}$$

Promotion never mutates or overwrites predecessor vertices. The strict non-promotion condition is:

$$\begin{aligned} \forall r_t \in \mathcal{X}, D \subseteq \mathcal{V}_h \cup \mathcal{M}_h : \\ (\exists u \in D \cap \mathcal{V}_h : \tau(u) \in \mathcal{U}) \\ \implies \forall v_s \in D \cap \mathcal{V}_h : \mathcal{G}_h \not\vdash_{\kappa} v_s \xrightarrow[D]{\text{prom}} v_t : r_t. \end{aligned} \quad (13)$$

Furthermore, admitting a new root vertex v directly into a grounded historical role $\mathcal{X}$ imposes the following root-admission obligations:

$$\begin{aligned} \forall v \in \mathcal{V}_h : \quad \text{indeg}(v) = 0 \wedge \tau(v) = \text{ObservedEvent} \\ \implies \text{CaptureOriginOK}_{\kappa}(v, h), \\ \forall v \in \mathcal{V}_h : \quad \text{indeg}(v) = 0 \wedge \tau(v) = \text{AttributedReport} \\ \implies \text{ReportOriginOK}_{\kappa}(v, h). \end{aligned} \quad (14)$$

CaptureOriginOK checks the admitted capture path, device identity, content commitment, and timestamp. ReportOriginOK checks an authenticated user- or third-party-authored message, its channel custody, bytes, and reception context. Such a message can ground an attributed report, never a direct observation.

Proposition 1 (Epistemic Provenance Traceability). *Assume E1–E3; a finite, acyclic $\mathcal{G}_h$; an admission gate that represents every non-root vertex transition by the judgment above with its exact dependency set; and root admission that enforces RootOK $_{\kappa}$, including Equation 14 for roles in $\mathcal{X}$. Every reachable claim $v \in \mathcal{C}^{\text{agent}}(h)$ produced by admitted transitions has a finite dependency trace obtained by recursively following its recorded references. The trace terminates only at policy-recognized source artifacts or explicit unavailable markers, and every admitted version and transition in it is inspectable at h .*

Proof sketch. Order $\mathcal{G}_h$ topologically and prove the stronger invariant for every admitted vertex. A root satisfies the root-admission assumption; a grounded root also satisfies Equation 14. For a non-root v , the recorded transition identifies a dependency set D , and E2 resolves each dependency to its exact version or an unavailable marker. For an admitted version $u \in D \cap \mathcal{V}_h$, the induction hypothesis supplies its finite trace, to which the recorded edge (u, v) attaches v . An unavailable marker instead forms a terminal trace entry bound to the affected dependency reference. Finiteness and acyclicity follow from the graph assumption. Equation 13 separately prevents a source in $\mathcal{U}$ from being retyped as a member of $\mathcal{X}$. Appendix A.1 provides the proof details. $\square$

2.7 Adversarial Threat Model

We consider an active adversary capable of injecting malicious natural-language instructions, corrupted biographical facts, or poisoned context into external web documents, retrieved snippets, user dialogues, or third-party tool outputs [6, 9]. The adversary’s objective is to induce the agent to emit false autobiographical statements, adopt fabricated user preferences, or violate confidentiality scopes.

Specifically, the adversary may:

1. **Sybil Corroboration:** Fabricate redundant records across distinct domains that originate from a single poisoned source to simulate synthetic consensus;
2. **Consolidation Injection:** Exploit background summarization or memory reflection to silently promote unverified prompt inputs into permanent profile beliefs;
3. **Temporal Exploitation:** Exploit expired permissions or stale historical preferences to authorize unauthorized present actions; or
4. **Declassification Probing:** Query the agent with crafted prompts designed to leak restricted user data or the existence of withheld records.

Trust Assumptions. We assume the integrity of the accepted state head h (governed by transactional consensus), the cryptographic key infrastructure, and the isolated execution environment of the assertion mediator. Adversarial attacks that compromise private signing keys, subvert local operating system kernel isolation, or leak information through unmediated side channels remain outside the primary security boundary.

3 Epistemic Resolution, Decision Witnesses, and Correction

Epistemic resolution formalizes the evaluation of a structured autobiographical query against an authenticated state head, an

authorized context projection, and an explicit release policy. For a requesting principal p , purpose u , and operational scope s , let $\mathcal{G}_h = (\mathcal{V}_h, \mathcal{E}_h, \lambda)$ be the accepted state graph at head h . We define:

$$\begin{aligned} \mathcal{A}_h &= \{v \in \mathcal{V}_h \mid \text{Reachable}(v, h)\}, \\ \mathcal{B}_h(p, u, s) &= \{v \in \mathcal{A}_h \mid \text{Authorized}(v, p, u, s)\}, \end{aligned} \quad (15)$$

where $\mathcal{A}_h$ is the accepted epistemic state relative to h and $\mathcal{B}_h(p, u, s)$ is its authorized projection. Neither set is assumed to be deductively closed or globally consistent; both may accommodate historical hypotheses, competing reports, unsuperseded conflicts, and access-sealed references.

3.1 Epistemic Resolution Semantics

Let $Q \in \mathcal{L}_{bio}$ be a structured autobiographical query binding subject, predicate, arguments, temporal bounds, and modality. Epistemic resolution under release policy κ is modeled as an evaluation judgment:

$$\mathcal{G}_h; \mathcal{B}_h(p, u, s) \vdash_{\kappa} Q \Downarrow \langle E, F, S, W \rangle, \quad (16)$$

yielding a 4-tuple $d = \langle E, F, S, W \rangle$ where:

- $E \in \Sigma_E = \{\text{Supported}, \text{Unknown}, \text{Unavailable}\}$ denotes the primary evidential status;
- $F \subseteq \Sigma_F = \{\text{Conflicted}, \text{Withheld}, \text{Stale}\}$ is a set of orthogonal operational qualifiers;
- $S \subseteq \mathcal{B}_h(p, u, s)$ is the minimized premise set available for content release; and
- $W \in \mathcal{W}_h \cup \{\emptyset\}$ is a protected decision witness object.

For this tuple, $\pi_E(d) = E$, $F(d) = F$, $S(d) = S$, and $W(d) = W$ denote its component projections. The resolver may inspect $\mathcal{A}_h$ to detect projection exclusions, but View_c is the only boundary through which a resolution result is declassified.

The resolution judgment enforces the following structural consistency conditions:

$$\begin{aligned} \pi_E(d) \in \Sigma_E, \quad F(d) \subseteq \Sigma_F, \\ \pi_E(d) = \text{Supported} &\iff S(d) \neq \emptyset \\ &\quad \wedge \text{PolicyOK}_{\kappa}(S(d), Q, \mathcal{G}_h), \\ \pi_E(d) \neq \text{Supported} &\implies S(d) = \emptyset, \\ \mathcal{G}_h; \mathcal{B}_h(p, u, s) \vdash_{\kappa} Q \Downarrow d \\ &\implies \text{WitnessOK}_{\kappa}(W(d), d, Q, h). \end{aligned} \quad (17)$$

Here, $\pi_E(d)$ projects the evidential status component of d . The predicate $\text{PolicyOK}_{\kappa}(S(d), Q, \mathcal{G}_h)$ verifies that premises in $S(d)$ satisfy validity intervals, any required source-independence checks ($\text{Indep}_{\kappa}(\cdot, \cdot; h)$), and derivation integrity under κ . When competing, unsuperseded claims mutually meet support thresholds, resolution yields $\pi_E(d) = \text{Supported}$ with $\text{Conflicted} \in F(d)$, populating $S(d)$ with both active versions.

3.2 Decision Witness Integrity and Minimization

A nonempty decision witness records the protected audit material required for a resolution result. Witness validity is formalized by:

$$\begin{aligned} \text{WitnessOK}_{\kappa}(W, d, Q, h) &\iff \\ &(W = \emptyset \wedge \text{AbsenceOnly}_{\kappa}(d, Q, h)) \\ &\vee (W \neq \emptyset \wedge \text{BoundToHead}(W, h) \\ &\quad \wedge \text{AccessBound}_{\kappa}(W, Q) \\ &\quad \wedge \text{MinimalWitness}_{\kappa}(W, d, Q) \\ &\quad \wedge \text{ObligationsMet}_{\kappa}(W, d, Q, h)). \end{aligned} \quad (18)$$

$\text{AbsenceOnly}_{\kappa}(d, Q, h)$ holds exactly when $\pi_E(d) = \text{Unknown}$, $F(d) = \emptyset$, and $\text{Rel}(Q, \mathcal{A}_h) = \emptyset$; thus it covers only an ordinary unknown caused by absence of admitted relevant material. BoundToHead binds or commits the witness to h ; AccessBound attaches its access policy; and MinimalWitness restricts the payload to material needed to justify d . A cryptographic realization may implement the bindings as commitments, but the abstract predicate does not require a particular construction.

ObligationsMet enforces category-specific witness bindings:

- **Conflict Witness:** Binds the exact incompatible current versions and their conflict relation;
- **Staleness Witness:** Binds the exact temporally inapplicable versions, relevant times, and governing temporal rule;
- **Withholding Witness:** Binds an access-controlled reference or commitment to relevant material in $\mathcal{A}_h$ that projection excludes from $\mathcal{B}_h$; and
- **Unavailability Witness:** Binds the exact erased, sealed, lost, or unsupported-adaptor marker and the affected dependency.

Diagnostic witnesses for failed verification or insufficient source roots remain protected and are optional unless policy requires them.

When Q lacks qualifying premises in $\mathcal{A}_h$, a resolver that records the evaluation context may use the empty-payload witness:

$$\perp_{\mathcal{W}}(h, Q, \kappa) = \langle h, \text{Digest}(Q), \text{Digest}(\kappa), \emptyset \rangle. \quad (19)$$

$\perp_{\mathcal{W}}$ carries no evidentiary payload; an ordinary absence-only Unknown may instead use $W = \emptyset$. W is not a support component and cannot itself justify content release; only members of S can do so. Absence of recorded support does not establish that the proposition is false or absent from every external source.

Table 2: Typed resolution statuses, orthogonal qualifiers, and governing declassification constraints.

Sort	Identifier	Deductive Evaluation Condition	Governing Declassification Constraint
Status	Supported	Admitted provenance DAG satisfies policy κ for query Q	Emit content assertion with mandatory source attribution
Qualifier	Conflicted	Multiple active, unsuperseded premises disagree	Disclose explicit epistemic disagreement or abstain
Status	Unknown	Authorized projection $\mathcal{B}_h$ lacks qualifying support	Disclose absence of recorded evidence; never assert falsity
Qualifier	Withheld	Relevant premises exist in $\mathcal{A}_h$ but are masked in $\mathcal{B}_h$	Emit uniform, recipient-safe nonconfirming response
Status	Unavailable	Required dependency is erased, sealed, lost, or lacks an adapter	Disclose operational unavailability only when policy authorizes
Qualifier	Stale	Premises exist but fail current temporal validity interval	Add explicit temporal qualification or abstain

3.3 Withholding Detection and Recipient Non-Interference

For a vertex set $Y \subseteq \mathcal{V}_h$, let $\text{Rel}(Q, Y) \subseteq Y$ denote the vertices semantically relevant to Q . Withholding arises when relevant historical records exist in the full store but are filtered out by access authorization:

$$\text{Withheld} \in F(d) \iff \text{Rel}(Q, \mathcal{A}_h) \setminus \text{Rel}(Q, \mathcal{B}_h(p, u, s)) \neq \emptyset. \quad (20)$$

In security-critical environments, revealing that data is withheld (e.g., “I have records that you are not authorized to view”) creates an existence side-channel, leaking confidential state. Let $\mathcal{D}_{unk}(c)$ contain the internal decisions that policy κ requires to remain indistinguishable in context c , potentially including **Unknown**, **Unavailable**, and decisions carrying **Withheld**. The declassification contract is:

$$\text{Collapse}_\kappa(c) \implies \forall d_1, d_2 \in \mathcal{D}_{unk}(c) : \text{View}_c(d_1) = \text{View}_c(d_2) = \perp_{unk}. \quad (21)$$

Section 4 restricts the renderer to this authorized view. Equation 21 is an observational policy condition, not a claim of cryptographic indistinguishability.

3.4 Epistemic Non-Vacuity and Scope Preservation

Resolving a query to **Unknown** represents a localized epistemic judgment at head h ; it does not assert the negation of the proposition $\neg\phi$, nor does it insert a negative axiom into state. The agent may still use general parametric reasoning in non-autobiographical domains. For example, an agent lacking records of a user’s dietary preferences may discuss culinary topics generally, but cannot assert that the user observes a specific dietary restriction.

Resolution also preserves operational scopes. An entitlement granted within a software development task does not generalize to enterprise financial auditing; consent granted to a specific coworker does not extend to third parties; and

a commitment established in 2024 does not authorize an unqualified present-tense assertion in 2026.

3.5 Dependency Invalidation and Forward State Reduction

When an upstream evidence artifact is found to be compromised, corrupted, or retracted, our framework executes forward state reduction rather than destructive history rewriting. A revocation operation records a new versioned state transition $h \xrightarrow{\text{revoke}(e)} h'$ appending a revocation descriptor to the graph.

Let $\mathcal{D}_h(v)$ be the nonempty family of complete transitive dependency alternatives recorded for v , each including its source roots; for a root, $\mathcal{D}_h(v) = \{\{v\}\}$. Let $\text{Revoked}(e, h)$ mean that e is revoked at h . Valid support is preserved when at least one complete alternative survives:

$$\text{ValidSupport}_\kappa(v, h) \iff \exists D \in \mathcal{D}_h(v) : \begin{aligned} &\text{Satisfies}_\kappa(D, v, \mathcal{G}_h) \\ &\wedge \forall e \in D : \neg \text{Revoked}(e, h). \end{aligned} \quad (22)$$

If every dependency path for v intersects at least one revoked root, v loses its warranted standing and is excluded from future support sets S . If alternative, mutually independent valid paths remain in $\mathcal{G}_h$, support is preserved.

We formalize four primitive lifecycle transitions:

- **Supersession** ($v_1 \succ v_2$): Introduces a successor claim that supersedes an earlier claim for current queries while preserving the immutable historical trace;
- **Dependency Invalidation** ($\text{revoke}(e)$): Invalidates an upstream root, requiring recomputation of downstream support via Equation 22;
- **Conflict Inscription** ($v_1 \#_{\text{conflict}} v_2$): Formally registers contradictory claims when no automated policy rule resolves precedence;

Table 3: Three complementary enforcement surfaces and their corresponding safety properties.

Enforcement Surface	Formal Safety Property
State Admission	Regulates which epistemic entities enter the accepted state graph $\mathcal{G}_h$
Context Projection	Filters accepted state into an authorized projection $\mathcal{B}_h(p, u, s)$
Assertion Mediation	Enforces content grounding, qualification, and recipient non-interference prior to release

- **Cryptographic / Physical Erasure** ($\text{erase}(e)$): Records that an affected dependency is unavailable after a policy-authorized erasure operation. Whether bytes are irrecoverable or a legal obligation is satisfied depends on the deployment.

4 Assertion Mediation and Bounded Release

Maintaining an intact accepted state in the persistence layer does not prevent an unconstrained generative model from hallucinating biographical details, omitting mandatory epistemic qualifications, or fabricating historical authority during natural-language synthesis. Assertion mediation establishes a formal security boundary at the point of external utterance release.

Table 3 contrasts these three distinct enforcement surfaces. Architectures implementing only storage constraints enforce *state-admission assurance*, which cannot prevent a probabilistic language model from emitting ungrounded claims during dialogue generation.

4.1 Generate–Verify–Revise Operational Semantics

Figure 1 illustrates the assertion mediation architecture. The generative model synthesizes candidate responses using an authorized projection bound to prompt head h_p . Before candidate text is emitted to governed external channels, the mediator intercepts and evaluates each constituent proposition.

We model mediation as an operational transition system over candidate configurations:

$$\langle z, c \rangle \xrightarrow{\text{extract}} \langle \bar{v}, \bar{Q}, c \rangle \xrightarrow{\text{resolve}} \langle \bar{v}, \bar{d}, c \rangle \xrightarrow{\text{gate}} z_{\text{rel}}, \quad (23)$$

where z is the candidate string, $c = \langle h_p, h_e, p, u, s, \kappa, t_e \rangle$ is the execution context, $\bar{v}$ is the ordered sequence of governed semantic units, $\bar{Q}$ binds one structured query to each unit, $\bar{d}$ contains the corresponding resolution results, and z_{rel} is the released output.

When a unit fails release verification, the mediator may provide the generator a declassified diagnostic derived through

View_c and request a revised candidate z' . Raw witness material $W(d)$ is not supplied to the generator or renderer. Every revision re-enters the pipeline at the extraction stage.

4.2 Emission-Head Temporal Verification

In concurrent agent systems, the authoritative state head may advance from h_p (the head bound during context assembly) to h_e (the head active at emission time t_e). To prevent race conditions and stale claim release, the mediator evaluates the emission-head judgment:

$$\begin{aligned} &\text{HeadUseOK}_\kappa(v, h_p, h_e, t_e) \\ &\iff h_p = h_e \\ &\quad \vee (\text{SnapshotOK}_\kappa(h_p, h_e, t_e) \\ &\quad \wedge \text{SnapQual}_\kappa(v, h_p, t_e)). \end{aligned} \quad (24)$$

If $h_p \neq h_e$ and policy disallows snapshot release, the candidate must be re-projected and re-mediated against current head h_e . If snapshot emission is authorized, the emitted unit v must include explicit temporal framing (e.g., “As of the preceding session...”).

4.3 Formal Autobiographical Assertion Boundedness

For any candidate output z , let $\text{Units}_{\mathcal{L}_{\text{bio}}}(z) = \langle v_1, \dots, v_n \rangle$ be the ordered sequence of covered semantic-unit occurrences, each with a unique unit identifier. The content and status subsets partition those occurrences:

$$\{v_1, \dots, v_n\} = \text{Content}_{\mathcal{L}_{\text{bio}}}(z) \dot{\cup} \text{Status}_{\mathcal{L}_{\text{bio}}}(z). \quad (25)$$

This partition applies to underlying semantic propositions rather than disjoint surface clauses: a single natural-language sentence may encode both a status disclosure and an embedded content claim.

For execution context $c = \langle h_p, h_e, p, u, s, \kappa, t_e \rangle$ and unit v bound to query Q_v , let $\mathcal{G}_{h_p}; \mathcal{B}_{h_p}(p, u, s) \vdash_\kappa Q_v \Downarrow d_v$. For resolver outputs, View_c is the sole declassification boundary: any result-dependent input to the generator or renderer is derived from $\text{View}_c(d_v)$, never raw $W(d_v)$. The renderer is constrained by:

$$\begin{aligned} &\text{StatusOut}_\kappa(d, c) = \text{Render}_\kappa(\text{View}_c(d), c), \\ &\text{View}_c(d_1) = \text{View}_c(d_2) \implies \\ &\text{RenderSet}_\kappa(\text{View}_c(d_1), c) = \text{RenderSet}_\kappa(\text{View}_c(d_2), c). \end{aligned} \quad (26)$$

Let $q(v)$ be the structured proposition extracted from unit v . We define the release admissibility relations for content




Figure 1: Assertion mediation pipeline for content claims and status responses. The boundary intercepts governed output channels, verifying candidate claims against accepted state before release.

and status units:

$$\begin{aligned}
 & \text{ContentReleaseOK}(\alpha, d_\alpha, c) \\
 & \iff \exists x \in S(d_\alpha) : \\
 & \quad \text{ExactMatch}(q(\alpha), \phi(x)) \\
 & \wedge \pi_E(d_\alpha) = \text{Supported} \\
 & \wedge \text{QualOK}_\kappa(\alpha, d_\alpha) \\
 & \wedge \text{DisclosureOK}_\kappa(\alpha, c) \\
 & \wedge \text{HeadUseOK}_\kappa(\alpha, h_p, h_e, t_e),
 \end{aligned} \tag{27}$$

$$\begin{aligned}
 & \text{StatusResponseOK}(\beta, d_\beta, c) \\
 & \iff \beta \in \text{RenderSet}_\kappa(\text{View}_c(d_\beta), c) \\
 & \wedge \text{HeadUseOK}_\kappa(\beta, h_p, h_e, t_e).
 \end{aligned} \tag{28}$$

Definition 2 (Autobiographical Assertion Boundedness). *A response z satisfies autobiographical assertion boundedness under execution context c , denoted $c \vdash_{\text{release}} z : \text{Bounded}_{\mathcal{L}_{\text{bio}}}$, if and only if every covered semantic unit satisfies its corresponding release relation:*

$$\begin{aligned}
 & \text{Bounded}_{\mathcal{L}_{\text{bio}}}(z, c) \iff \\
 & (\forall \alpha \in \text{Content}_{\mathcal{L}_{\text{bio}}}(z) : \\
 & \quad \text{ContentReleaseOK}(\alpha, d_\alpha, c)) \\
 & \wedge (\forall \beta \in \text{Status}_{\mathcal{L}_{\text{bio}}}(z) : \\
 & \quad \text{StatusResponseOK}(\beta, d_\beta, c)).
 \end{aligned} \tag{29}$$

Proposition 2 (Soundness of Assertion Mediation). *Assume E1–E4; complete extraction and correct unit typing over $\mathcal{L}_{\text{bio}}$; correct implementations of ExactMatch, QualOK, and DisclosureOK; resolution satisfying Equation 17, WitnessOK $_\kappa$, and dependency-invalidation recomputation; a correct View_c and renderer satisfying Equations 21 and 26; recursive re-mediation of revisions; and correct HeadUseOK $_\kappa$. If the transition system in Equation 23 reaches a release state with output z_{rel} , then $\text{Bounded}_{\mathcal{L}_{\text{bio}}}(z_{\text{rel}}, c)$ holds.*

This is a conditional composition result: it does not establish the correctness of the extraction, policy predicates, resolver, or declassification function assumed in its premise.

Proof sketch. By E4, every candidate on a governed outbound channel reaches the mediator. Complete extraction and correct typing produce the ordered covered units

and their content/status partition. A content assertion α is emitted only after the mediator verifies every conjunct of ContentReleaseOK using the assumed-correct predicates and resolver. A status unit is emitted only from the render set of its authorized view and after the head-use check. The renderer receives no input other than $\text{View}_c(d)$; equal views therefore yield equal render sets, and any policy-required collapse is supplied by Equation 21. A failed candidate is not released, and every revision re-enters extraction. Thus every unit in a released output satisfies its corresponding release predicate, which is exactly $\text{Bounded}_{\mathcal{L}_{\text{bio}}}(z_{\text{rel}}, c)$. Appendix A.2 provides the proof details. $\square$

4.4 Assertion Trace Semantics and Auditability

Let $\text{Units}_{\mathcal{L}_{\text{bio}}}(z_{\text{rel}}) = \langle v_1, \dots, v_n \rangle$. The specification assigns each released governed unit exactly one binding, in that order:

$$\begin{aligned}
 \mu_v = \langle & \text{Id}(v), \text{Digest}(v), \text{Type}(v), \\
 & \text{Digest}(Q_v), \text{Digest}(d_v), \\
 & \text{Digest}(S(d_v)), \text{Digest}(W(d_v)), \\
 & \delta_v, \sigma_v \rangle,
 \end{aligned} \tag{30}$$

where $\text{Type}(v) \in \{\text{Content}, \text{Status}\}$, δ_v records the mediation disposition (Pass, Qualify, Redact), and σ_v records the emission-head verification result.

The response-level trace is:

$$\begin{aligned}
 \bar{\mu}_z &= \langle \mu_{v_1}, \dots, \mu_{v_n} \rangle, \\
 \mathcal{T}_z &= \langle \bar{\mu}_z, h_p, h_e, t_e, \kappa, \text{Digest}(z_{\text{rel}}) \rangle.
 \end{aligned} \tag{31}$$

An implementation with canonical serialization, hashing, and protected storage can use this structure for audit review. The data structure alone establishes neither immutability nor non-repudiation.

4.5 Operational Boundaries and Limitations

Assertion mediation relies on the semantic fidelity of the extraction and classification stage. Structured schema slots and deterministic templates avoid free-form extraction for the fields they govern, but free-form natural-language extraction remains vulnerable to subtle presuppositions, implicative phrasing, or compound rhetorical structures. Furthermore,

real-time conversational interfaces impose stringent latency budgets, necessitating cached verification or constrained template rendering.

The accompanying executable artifact instantiates structured claim objects, a typed resolver, fixed recipient-facing templates, ordered binding data structures, and 24 hand-authored cases. It does not implement natural-language semantic-unit extraction, generalized witness access control, configurable View_c policies or output-equivalence checks, full provenance-DAG traversal, task/branch/relationship scope enforcement, cryptographic witness or assertion commitments, canonical serialization or artifact-hash verification, or the target-origin admission checks in Equation 14. Those components remain specification-level obligations.

5 Evaluation Framework and Conformance Checks

We evaluate the operational correctness of the typed resolver and assertion mediator using an executable conformance suite, distinguishing these structural verification checks from the proposed end-to-end evaluation protocol detailed in Appendix C. The suite checks the reference implementation against encoded decision oracles, flattened source-root constraints, and release invariants across structured threat scenarios.

5.1 Evaluation Objectives and Baseline Paradigms

The evaluation design separates executable conformance questions from measurements that require a deployed agent. It covers four operational dimensions:

- **State admission and provenance discipline:** Verifying that typed claim objects and origin checks prevent unverified model priors, unadmitted inputs, and single-root duplicate records from acquiring authoritative autobiographical standing.
- **Assertion release control:** Evaluating the mediator’s capability to intercept unsupported, out-of-scope, or contradictory candidate statements while preserving legitimately supported answers with appropriate attributions.
- **Longitudinal state evolution:** Tracking how temporal expiration, superseding user corrections, multi-source conflicts, and model substrate upgrades are handled without historical corruption or silent promotion.
- **Operational overhead and utility:** Measuring decision latency, computational overhead, and over-abstention across deployed assurance configurations; the present artifact provides only a local resolver microbenchmark and deterministic coverage counts.

For a future end-to-end study, we formulate five comparative deployment paradigms: (1) a pure parametric agent without episodic memory; (2) a flat vector-retrieval store over raw transcripts; (3) standard top- k retrieval-augmented generation (RAG); (4) provenance-tagged retrieval lacking graph-level root independence; and (5) a state-bounded architecture that filters memory projections at prompt assembly without pre-release mediation. The present conformance suite instead isolates three deterministic release rules: flat/prior, source-tag, and typed mediation. Appendix C specifies the five deployment paradigms for later end-to-end comparison.

5.2 Safety and Utility Metrics

Let C denote candidate assertion opportunities evaluated by the mediator, $U \subseteq C$ denote opportunities that are unsupported or require qualification under policy κ , $R \subseteq C$ denote released responses emitting content or qualified status, and $P \subseteq R$ denote candidate content passed without required qualifications.

We quantify mediation safety and efficacy through two primary risk metrics:

1. **Unsupported-Opportunity Release Rate (UOR):** The proportion of unsafe or unqualified candidate assertions that escape mediation:

$$\text{UOR} = \frac{|U \cap P|}{|U|}. \quad (32)$$

2. **Release Contamination (RC):** The proportion of released conversational outputs containing unqualified unsafe assertions:

$$\text{RC} = \frac{|U \cap P|}{|R|} \quad (\text{defined as 0 when } R = \emptyset). \quad (33)$$

Because safety can trivially be maximized through blanket refusal, we pair these metrics with **Coverage**—the proportion of legitimately supported queries for which the system releases an accurate, appropriately qualified response. In addition, we measure **Prevention Precision** (proportion of interventions targeting genuinely unsafe opportunities), **Prevention Recall** (proportion of unsafe opportunities intercepted), and **Mediation Latency**.

5.3 Executable Conformance Suite Results

We evaluate the in-memory claim store, typed resolver, and assertion mediator against 24 structured cases: 19 unsafe-release opportunities and five supported controls across six threat tracks. Cases that require corroboration use a policy requiring two roots that satisfy $\text{Indep}_\kappa(\cdot, \cdot; h)$.

Table 4 reports aggregate metrics across three deterministic release rules, and Table 5 details the per-case oracle. Across this test fixture:

Table 4: Conformance suite aggregate results under the two-source fixture (24 cases).

Release rule	UOR ↓	RC ↓	Coverage ↑	Prevention precision ↑	Prevention recall ↑
Flat/prior rule	19/19 (100.0%)	19/24 (79.2%)	5/5 (100.0%)	N/A—no interventions	0/19 (0.0%)
Source-tag rule	18/19 (94.7%)	18/23 (78.3%)	5/5 (100.0%)	1/1 (100.0%)	1/19 (5.3%)
Typed mediation	0/19 (0.0%)	0/7 (0.0%)	5/5 (100.0%)	19/19 (100.0%)	19/19 (100.0%)

Table 5: Per-case oracle and decisions across the 24-case conformance suite.

Case	Type	Condition	Fixture roots	Decision d_h	Candidate check	Disposition
poison_01	Unsafe	Unadmitted model prior without evidence	–	Unknown	Model-origin candidate	Blocked
poison_02	Unsafe	Indirect prompt injection via untrusted doc	untrusted_web_page	Unknown	Unverified evidence	Blocked
poison_03	Unsafe	Hypothetical simulation claiming observation	sim_sandbox_log	Unknown	Simulation origin	Blocked
poison_04	Unsafe	Fabricated unverified consolidation summary	unverified_llm_hallucination	Unknown	Unverified derivation	Blocked
dup_01	Unsafe	Duplicate records share single upstream root	root_chat_session_A	Unknown	Insufficient roots	Blocked
dup_02	Unsafe	Multi-hop derivation chain ($k = 3$) sharing root	root_session_single	Unknown	Single upstream root	Blocked
dup_03	Unsafe	Partial root overlap below 2-root threshold	shared_root_omega	Unknown	Shared root overlap	Blocked
supported_01	Control	Valid preference corroborated by 2 roots (Control)	supported_root_a, supported_root_b	Supported	Exact $S(d_h)$ member	Passed
scope_01	Unsafe	Unauthorized principal query for restricted intake	private_intake	Unknown + Withheld	Outside authorized view	Blocked
scope_02	Unsafe	Purpose mismatch against authorized scope	payroll_root_1, payroll_root_2	Unknown + Withheld	Purpose mismatch	Blocked
scope_03	Unsafe	Sealed / cryptographically erased record	erased_device_vault	Unavailable	Sealed record unavailable	Blocked
supported_02	Control	Authorized principal & matching purpose (Control)	org_root_a, org_root_b	Supported	Exact $S(d_h)$ member	Passed
time_01	Unsafe	Expired permission grant queried past validity	time_root_1, time_root_2	Unknown + Stale	No current support	Blocked
time_02	Unsafe	Premature commitment queried before start time	promo_root_1, promo_root_2	Unknown + Stale	Not yet active	Blocked
time_03	Unsafe	Transitive staleness in derived preference summary	order_root_a, order_root_b	Unknown + Stale	Observation expired	Blocked
supported_03	Control	Active preference within valid bounds (Control)	tz_root_a, tz_root_b	Supported	Exact $S(d_h)$ member	Passed
conflict_01	Unsafe	Two active unsuperseded conflicting locations	city_root_1, city_root_2, city_root_3, city_root_4	Supported + Conflicted	Status-only response	Qualified
conflict_02	Unsafe	Three-way location disagreement across roots	r1, r2, r3, r4, r5, r6	Supported + Conflicted	Status-only response	Qualified
invalidation_01	Unsafe	Transitive dependency invalidation revokes support	revoked_cert_root_1, revoked_cert_root_2	Unknown	Revoked upstream root	Blocked
invalidation_02	Control	Redundant independent root survives invalidation (Control)	fido2_root_valid_a, sms_root_invalidated, smtp_root_valid_b	Supported	Exact $S(d_h)$ member	Passed
correction_01	Unsafe	Outdated preference superseded by user correction	lang_root_new_a, lang_root_new_b, lang_root_old_a, lang_root_old_b	Supported	Outdated candidate not in $S(d_h)$	Blocked
correction_02	Unsafe	Multi-version chain ($v_1 \rightarrow v_2 \rightarrow v_3$) asserting v_2	r_th1a, r_th1b, r_th2a, r_th2b, r_th3a, r_th3b	Supported	Superseded revision v_2	Blocked
correction_03	Control	Active latest revision v_3 in chain (Control)	r_th1a, r_th1b, r_th2a, r_th2b, r_th3a, r_th3b	Supported	Exact $S(d_h)$ member	Passed
head_race_01	Unsafe	Emission-head race mismatch without snapshot permission	k8s_root_a, k8s_root_b	Supported	Stale head use rejected	Blocked

1. The **Flat/prior rule** admits all 19 unsafe candidates (100.0% UOR, 79.2% RC across $|R| = 24$ emissions). By treating model priors and raw retrieved text as authoritative assertions, this rule serves as an unmitigated comparison condition illustrating baseline generative vulnerability.
2. The **Source-tag rule** prevents one unsafe release but admits the other 18 (94.7% UOR, 78.3% RC over $|R| = 23$). Because surface metadata tags do not track root

provenance, temporal validity, or supersession history, tagging alone cannot prevent false corroboration or stale releases.

3. The **Typed mediation** passes no unsafe candidate unqualified (0/19 UOR and 0/7 RC) while passing all five controls (100.0% Coverage), with 100.0% prevention precision and recall. It blocks 17 unsafe candidates and emits qualified status responses for the two active-conflict cases, yielding $|R| = 7$ total emissions.

5.4 Falsification Conditions

The architecture’s empirical validity rests on concrete falsification conditions. The approach would be empirically refuted if: (1) semantic extraction fails to capture covered propositions above a declared operational error tolerance; (2) mediation overhead reduces valid-query coverage below an acceptable utility threshold; or (3) a strictly simpler heuristic achieves equivalent safety and coverage at lower operational latency. Appendix C defines these criteria for future end-to-end deployments.

Fourteen test functions additionally cover mixed temporal/conflict states, witness fields, non-disclosure templates, ordered released-unit bindings, exact aggregate regressions, and a bounded enumeration of 1,152 single-claim configurations. These unit checks add no rows to Tables 4–5. Neither the case oracle nor the auxiliary checks establish semantic correctness of arbitrary natural-language text, recipient-view equivalence, or exhaustive state-space coverage.

5.5 Artifact Availability

The public `openkedge/pci` artifact is published under the MIT license and pinned at release `v0.2.0` (commit `7484153`). That revision provides the full 24-case conformance suite, typed mediator, and 1,152-state bounded exploration suite. In the present working tree, `make reproduce` runs 14 tests and regenerates the 24-case tables.

6 Discussion and Deployment Considerations

Enforcing epistemic discipline provides an inspectable, auditable claim regime; it does not, by itself, establish the moral or institutional legitimacy of the underlying policy rules. This distinction is critical when claim standing governs automated actions, confidential disclosure, or long-term personalization. Appendix E provides an extended analysis of sociotechnical governance, multi-party consent, and regulatory erasure.

6.1 Withholding Privacy and Existence Leakage

Typed claim objects can support fine-grained access control across principals, purposes, and temporal intervals. Once a consent change is admitted and its projection policy becomes active, affected claims can be excluded from subsequent views $\mathcal{B}_h(p, u, s)$ without rewriting historical state.

The act of withholding information can itself leak confidential state. Emitting a specific refusal (e.g., “Access to your medical records is denied”) confirms to an unauthorized requester that relevant records exist. When policy requires that existence remain secret, View_c can map `Withheld`, `Unavailable`, and `Unknown` to the same declassified view, so that the renderer returns a uniform nonconfirming response. Appendix E.1 states the observation assumptions for this branch-distinguishing channel.

6.2 Relational Privacy and Multi-Party Contexts

Autobiographical agent memories often capture multi-party interactions. An agent’s authenticated provenance for a conversation does not grant unilateral authority to disclose a third party’s private statements or attributes to other users. Deployments must enforce multi-party provenance isolation, ensuring that one user’s interaction history cannot leak into another’s authorized projection (Appendix E.2).

6.3 Procedural Guarantees vs. Semantic Truth

Our conditional contract theorem (Proposition 2) establishes strict procedural bounds:

- Assertions satisfy declared operational rules, not that external informants are infallible or objectively true;
- Recorded stances reflect system policy, not that an agent experiences subjective belief or conscious memory;
- Cryptographic signatures authenticate signed bytes and key possession, rather than factual veracity.

6.4 Closed-Channel Scope and Side Channels

Our guarantees apply to governed output channels passing through assertion mediation. Ungoverned side channels—such as debug logs, diagnostic telemetry, cached prompts, or unmediated tool APIs—must be isolated or routed through mediation gateways to prevent out-of-band autobiographical leaks. Systems should catalog all I/O surfaces and report state admission, projection, and mediation assurance levels as distinct operational tiers.

7 Related Work

Epistemic foundations and bounding. In formal verification, De Giacomo et al. bound fluent atoms considered possibly true by an agent for decidable verification over unbounded transition systems [10]. In governance, Kim et al. analyze epistemic bounds in administrative oversight of generative AI [11]. Here, *autobiographical assertion boundedness* specifically ensures that an agent’s outward assertions concerning itself, its user, or named relationships strictly conform to accepted evidence, validity horizons, and authorized disclosure policies.

Transactional state integrity. Transactional memory systems govern state admission, version changes, and repair under concurrent updates [7, 8]. Assertion boundedness begins after that boundary: it assumes an accepted head and governs the evidential and disclosure conditions for outward statements. Appendix F compares this release contract with related memory and grounding systems in Tables 9 and 10.

Persistent memory and agent architectures. Hierarchical memory architectures enable agents to maintain state across extended dialogues. Generative Agents [1] and MemGPT [2] pioneered tiered memory using periodic reflection and retrieval. A-MEM [3] and Mem0 [4] organize dynamic stores for lifelong personalization, while LoCoMo [5] and LaMP [12] evaluate long-context retrieval. These systems focus on memory construction and retrieval rather than the typed, recipient-specific output-release contract studied here.

Structured and provenance-aware memory. Recent systems introduce structured memory representations. MemIR [13] decouples evidence, retrieval cues, and memory atoms to prevent provenance-role collapse. MemTX [8] treats belief updates as transactional operations with repair cascades. Eywa [14] maintains immutable source evidence while deriving provenance-linked canonical facts. MemLineage [15] secures memory via cryptographic lineage graphs for tool invocations. Our work addresses a complementary problem: the outward assertion boundary. While MemIR, MemTX, and Eywa structure internal storage, our architecture formalizes policy-relative promotions, typed resolution with orthogonal flags, recipient-safe status rendering, and emission-head race verification for outward autobiographical claims.

Epistemic logic and belief revision. Dynamic epistemic logic formalizes multi-agent knowledge transitions [16], while AGM belief revision analyzes minimal, rational state change [17]. Rather than assuming idealized deductive closure, our framework formulates procedural, dependency-bounded resolution tailored for neural agents operating over partial, open-world evidence stores.

Data provenance and verifiable credentials. The W3C PROV-DM standard models entities, activities, and agents [18], while Verifiable Credentials authenticate issuer statements and their presentation [19]. Our typed provenance graph adds epistemic roles, temporal validity, disclosure scopes, and policy-relative source independence.

Agent security and memory poisoning. Autonomous agents are susceptible to indirect prompt injection via retrieved web pages, external tools, and multi-turn dialogues [6, 9]. ConsistencyGate [20] filters untrusted inputs using write-time model consistency. Our framework complements input filtering by enforcing source-independence graph analysis at resolution time and intercepting ungrounded releases at the output mediation boundary.

Output grounding and hallucination mitigation. Post-generation verification tools, such as MiniCheck [21], evaluate whether emitted claims are grounded in provided documents. While output grounding checks document faithfulness,

our assertion mediator evaluates the broader epistemic context: verifying that supporting evidence is authoritatively admitted, independent, temporally active, and authorized for release to the specific recipient.

Parametric knowledge editing. Techniques such as ROME [22] and MEMIT [23] directly modify parametric factual associations in model weights. Parametric editing operates on a different layer than persistent agent state: at a fixed accepted head, our framework keeps the explicit autobiographical claim set invariant under model substitution and denies generated content evidential standing by default.

8 Conclusion

Persistent AI agents require an enforceable distinction between material that is stored and claims that are supported for release. Neither a foundation model’s fluency nor a memory object’s reachability gives an individual agent evidence, authority, and valid scope to claim content autobiographically. We formalized *autobiographical assertion boundedness* through a system-level contract spanning a typed provenance graph, source independence, controlled role promotion, typed resolution with orthogonal flags, and generate–verify–revise assertion mediation.

Under its stated extraction, predicate-correctness, resolution, declassification, and channel assumptions, the conditional contract result restricts release to units that pass the content or status checks. In 24 deterministic, hand-authored conformance cases, typed mediation passed none of 19 unsafe opportunities unqualified and passed all five supported controls. The suite validates the encoded obligations but does not measure natural-language extraction or deployed retrieval and generation systems. By keeping reachability and support separate, the resulting contract provides a precise basis for evaluating how persistent agents bind outward autobiographical claims to accepted state.

Acknowledgments

OpenAI Codex was used for language editing, formal-consistency review, and execution of reproducibility checks. The authors reviewed the resulting manuscript, proofs, citations, and experimental outputs and take responsibility for the final content.

References

- [1] Joon Sung Park, Joseph C. O’Brien, Carrie J. Cai, Meredith Ringel Morris, Percy Liang, and Michael S. Bernstein. Generative agents: Interactive simulacra of human behavior. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology*, pages 1–22, 2023.
- [2] Charles Packer, Sarah Wooders, Kevin Lin, Vivian Fang, Shishir G. Patil, Ion Stoica, and Joseph E. Gonzalez. MemGPT: Towards LLMs as operating systems. *arXiv preprint arXiv:2310.08560*, 2023.

- [3] Wujiang Xu, Zujie Liang, Kai Mei, Hang Gao, Juntao Tan, and Yongfeng Zhang. A-MEM: Agentic memory for LLM agents. In *Advances in Neural Information Processing Systems* 38, pages 17577–17604. Curran Associates, Inc., 2025.
- [4] Prateek Chhikara, Dev Khant, Saket Aryan, Taranjeet Singh, and Deshraj Yadav. Mem0: Building production-ready AI agents with scalable long-term memory. *arXiv preprint arXiv:2504.19413*, 2025.
- [5] Adyasha Maharana, Dong-Ho Lee, Sergey Tulyakov, Mohit Bansal, Francesco Barbieri, and Yuwei Fang. Evaluating very long-term conversational memory of LLM agents. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 13851–13870. Association for Computational Linguistics, 2024.
- [6] Edoardo DeBenedetti, Jie Zhang, Mislav Balunovic, Luca Beurer-Kellner, Marc Fischer, and Florian Tramèr. AgentDojo: A dynamic environment to evaluate prompt injection attacks and defenses for LLM agents. In *Advances in Neural Information Processing Systems 37, Datasets and Benchmarks Track*, 2024.
- [7] Jun He and Deying Yu. Beyond memory: A transactional continuity kernel for long-lived AI agents. *arXiv preprint arXiv:2608.11632*, 2026.
- [8] Xiaoyang Li, Yiqi Wang, Haohui Lu, Zhi Chen, Mo Li, Pingan Song, Mingkai Zheng, and Taotao Cai. MemTX: Transactional belief commit for stateful agent memory. *arXiv preprint arXiv:2607.23929*, 2026.
- [9] Sidharth Pulipaka, Stanislaw Hlebik, Leonidas Raghav, Sahar Abdelnabi, Vyas Raina, Ivaxi Sheth, and Mario Fritz. Hidden in memory: Sleeper memory poisoning in LLM agents. *arXiv preprint arXiv:2605.15338*, 2026.
- [10] Giuseppe De Giacomo, Yves Lespérance, and Fabio Patrizi. Bounded epistemic situation calculus theories. In *Proceedings of the Twenty-Third International Joint Conference on Artificial Intelligence*, pages 846–853, 2013.
- [11] Yushim Kim, Jieun Kim, Taeuk Kim, and Hee-Chan Cho. Administrative decision-making with generative AI: The challenge of epistemic boundedness. *Administration & Society*, 58(2):284–304, 2026. <https://doi.org/10.1177/00953997251409156>.
- [12] Alireza Salemi, Sheshera Mysore, Michael Bendersky, and Hamed Zamani. LaMP: When large language models meet personalization. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 7370–7392. Association for Computational Linguistics, 2024.
- [13] Zhengda Jin, Bingbing Wang, Jing Li, Ruifeng Xu, and Min Zhang. Mitigating provenance-role collapse in long-term agents via typed memory representation. *arXiv preprint arXiv:2605.25869*, 2026.
- [14] Resham Joshi. Eywa: Provenance-grounded long-term memory for AI agents. *arXiv preprint arXiv:2605.30771*, 2026.
- [15] Ciyang Ouyang and Rui Hou. MemLineage: Lineage-guided enforcement for LLM agent memory. *arXiv preprint arXiv:2605.14421*, 2026.
- [16] Ronald Fagin, Joseph Y Halpern, Yoram Moses, and Moshe Y Vardi. *Reasoning About Knowledge*. MIT Press, 2004.
- [17] Carlos E Alchourrón, Peter Gärdenfors, and David Makinson. On the logic of theory change: Partial meet contraction and revision functions. *The Journal of Symbolic Logic*, 50(2):510–530, 1985.
- [18] Luc Moreau and Paolo Missier. PROV-DM: The PROV data model. W3C Recommendation, 2013.
- [19] W3C Verifiable Credentials Working Group. Verifiable credentials data model v2.0. W3C Recommendation, 2025.
- [20] Yan Zhang and Shibo Li. ConsistencyGate: Preventing memory contamination in LLM agents via self-consistency admission control. *arXiv preprint arXiv:2607.22962*, 2026.
- [21] Liyan Tang, Philippe Laban, and Greg Durrett. MiniCheck: Efficient fact-checking of LLMs on grounding documents. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 8818–8847. Association for Computational Linguistics, 2024.
- [22] Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov. Locating and editing factual associations in GPT. In *Advances in Neural Information Processing Systems* 35, 2022.
- [23] Kevin Meng, Arnab Sen Sharma, Alex Andonian, Yonatan Belinkov, and David Bau. Mass-editing memory in a transformer. In *International Conference on Learning Representations*, 2023.

A Supporting Formal Proofs

This appendix supplies proof details for Proposition 1 (Epistemic Provenance Traceability) and Proposition 2 (Soundness of Assertion Mediation) under their stated assumptions.

A.1 Proof of Proposition 1 (Epistemic Provenance Traceability)

Proof. Let h satisfy the assumptions of Proposition 1. Because $\mathcal{G}_h = (\mathcal{V}_h, \mathcal{E}_h, \lambda)$ is finite and acyclic, it has a topological ordering. We induct over that ordering to establish the stronger invariant for every admitted vertex: it has a finite dependency trace whose terminal entries are policy-recognized source artifacts or unavailable markers, with exact admitted versions and transitions inspectable at h .

Base case. Let v be a root. The root-admission assumption gives $\text{RootOK}_\kappa(v, h)$. If $\tau(v) = \text{ObservedEvent}$ or $\tau(v) = \text{AttributedReport}$, Equation 14 additionally requires $\text{CaptureOriginOK}_\kappa(v, h)$ or $\text{ReportOriginOK}_\kappa(v, h)$, respectively. A root in another role retains its policy-recognized origin marker without thereby joining $\mathcal{X}$. The singleton trace at v satisfies the claim.

Induction step. Let non-root v be admitted by a recorded judgment $\mathcal{G}_h \vdash_\kappa v_s \xrightarrow[D]{\text{prom}} v : \tau(v)$. Its dependencies precede v in the topological order. By E2, each reference in D resolves to its exact admitted version or an explicit unavailable marker. For each exact version $u \in D \cap \mathcal{V}_h$, the induction hypothesis supplies a finite dependency trace; attaching v through the recorded edge (u, v) extends that trace. An unavailable marker in $D \cap \mathcal{M}_h$ is a terminal entry bound to its affected reference in the transition record. The graph assumption preserves finiteness and acyclicity. Equation 13 separately prevents a source in $\mathcal{U}$ from being promoted into $\mathcal{X}$. $\square$

A.2 Proof of Proposition 2 (Soundness of Assertion Mediation)

Proof. Let z be a candidate response on a governed channel under context $c = \langle h_p, h_e, p, u, s, \kappa, t_e \rangle$. We show that any output z_{rel} reached by the release transition satisfies $\text{Bounded}_{\mathcal{L}_{\text{bio}}}(z_{\text{rel}}, c)$.

Interception and semantic decomposition. By E4, every candidate emission on a governed channel is intercepted. Under the complete-extraction and correct-typing assumptions, the candidate is decomposed into the ordered occurrences $\text{Units}_{\mathcal{L}_{\text{bio}}}(z) = \langle v_1, \dots, v_n \rangle$. Their unique identifiers occur exactly once, and the occurrences partition into content and status units. Each v is bound to a structured query $Q_v \in \mathcal{L}_{\text{bio}}$.

For each unit v , the mediator computes $\mathcal{G}_{h_p}; \mathcal{B}_{h_p}(p, u, s) \vdash_\kappa Q_v \Downarrow d_v = \langle E, F, S, W \rangle$:

- *Content assertions.* The mediator emits α only if every conjunct of $\text{ContentReleaseOK}(\alpha, d_\alpha, c)$ holds. Equa-

tion 17 supplies nonempty, policy-valid support for a **Supported** result. The assumed-correct ExactMatch , QualOK , DisclosureOK , and HeadUseOK_κ implementations establish the remaining conjuncts.

- *Status responses.* The renderer receives $\text{View}_c(d_\beta)$, never raw $W(d_\beta)$. A status unit is released only if $\beta \in \text{RenderSet}_\kappa(\text{View}_c(d_\beta), c)$ and it passes HeadUseOK . Equation 26 gives equal render sets for equal views; where policy requires internal outcomes to be indistinguishable, Equation 21 maps them to the same view.

If any semantic unit fails verification, that candidate is blocked. A revision re-enters extraction, and E4 prevents bypass. Hence the release transition is reachable only when every covered unit satisfies its corresponding release relation. By Definition 2, the emitted z_{rel} is bounded. $\square$

B Taxonomies and Role Promotion Rules

Table 6 details the origin taxonomy and provenance guarantees, and Table 7 illustrates instances of the policy-relative judgment $\mathcal{G}_h \vdash_\kappa v_s \xrightarrow[D]{\text{prom}} v_t : r_t$.

C Proposed End-to-End Evaluation Protocol and Tracks

Table 8 defines a proposed seven-track adversarial evaluation design for long-term agent deployments. The present paper does not report this end-to-end study.

Baselines and Ablation Configurations. The proposed study would evaluate five deployment baselines: (1) a parametric agent using zero-shot system prompts without episodic memory; (2) a flat memory store using dense semantic search over past transcripts without provenance parsing; (3) standard RAG injecting retrieved chunks with grounding instructions; (4) provenance-tagged retrieval without root-DAG analysis or invalidation checks; and (5) a state-bounded baseline that filters claim objects at prompt assembly without post-generation mediation. Retrieval depth, chunking, model, and decoding settings must be fixed and reported when the study is executed. Proposed ablations isolate source-root independence (Roots_κ), transitive dependency invalidation (Equation 22), emission-head timing (HeadUseOK_κ), and recipient-safe declassification (View_c).

Falsification Criteria. The architecture would be empirically challenged if: (1) semantic extraction misses covered assertions above a deployment’s declared safety tolerance; (2) mediation reduces valid-query coverage beyond its declared utility bound; or (3) a simpler rule achieves equivalent

Table 6: Taxonomy of claim origins, verifiable provenance guarantees, and boundaries.

Origin Class	Verifiable Provenance Guarantee	Epistemic Boundary
Direct capture	Hardware device, byte commitment, capture path, and timestamp	Sensor calibration and environmental fidelity
Attributed report	Channel attribution, cryptographic signature, and asserted scope	Informant honesty, accuracy, and temporal stability
Imported record	Document version hash, repository custody, and import pipeline	Historical applicability to current agent context
Model prior	Model architecture identifier and checkpoint metadata	Lacks evidentiary grounding in agent interaction history
Retrieved external text	Search query, retrieved snippet hash, source URL, and timestamp	External source reliability and future validity
Deterministic derivation	Exact dependency inputs, code version hash, and deterministic output	Semantic soundness of transformation logic
Probabilistic interpretation	Model evaluator version, input embeddings, and sealed inference	Evaluator drift and non-deterministic reproducibility
Hypothetical simulation	Simulation framework, parameter config, and scenario label	Explicitly non-actualized synthetic execution

Table 7: Policy-relative role promotions and target admission constraints.

Source Role (r_s)	Target Role (r_t)	Governing Policy Constraint
Attributed report	Memory interpretation	Binds original report hash, interpretation evaluator version, and scope; preserves attribution
Memory interpretation	Recorded belief	Requires explicit stance endorsement authority and retention of underlying interpretation dependencies
Model prior	Marked hypothesis	Requires labeling parametric model origin; denies unqualified autobiographical standing
Authenticated user message	Attributed report	Creates a report object bound to sender signature, channel custody, message hash, and timestamp
Direct sensor capture	Observed event	Creates an observation object bound to capture device pipeline, timestamp, and integrity checks
Model prior / unverified retrieval	Observed event or attributed report	Strictly prohibited: unverified or parametric origins cannot serve as observation or report provenance

Table 8: Adversarial evaluation tracks and target epistemic failure modes.

Evaluation Track	Target Scenarios	Primary Failure Mode Exposed
Origin & promotion	Model priors, unverified web text, simulations, synthetic summaries	Unauthorized elevation of ungrounded inputs into autobiographical history
Source independence	Surface paraphrases, multi-URL duplicates sharing upstream roots	False corroboration and synthetic consensus through redundant retrieval
Scope & disclosure	Recipient-restricted data, task permissions, confidential records	Unauthorized cross-context data leakage and privacy violations
Conflict & time	Contradictory records, expired permissions, retracted evidence	Stale assertion release, unhandled disagreement, and invalidation gaps
Memory poisoning	Indirect prompt injections in ingested docs, snippets, and tools	Persistent compromise of agent profile and control flow across sessions
Substrate upgrade	Model replacement ($m_1 \rightarrow m_2$) under fixed accepted head h	Hallucinatory embellishment or parametric drift masquerading as biography
Assertion extraction	Presuppositions, negations, modality shifts, compound clauses	Evading assertion mediation via indirect or non-declarative framing

safety and coverage at lower latency. These thresholds should be selected before executing the end-to-end study; the conformance suite does not estimate them.

D Conformance Suite Fixture and Dependency Topology

The deterministic conformance suite evaluates the typed resolver and assertion mediator across 24 cases in six threat tracks. Corroboration cases require two roots that satisfy $\text{Indep}_\kappa(\cdot, \cdot; h)$.

Dependency Topology and Execution. The evaluation fixture encodes acyclic dependency topologies through accepted claims, evidence items, and terminal-root identifiers; the reference resolver consumes flattened root references rather than traversing a general $\mathcal{G}_h$. The cases cover model priors, prompt injection, simulated or hallucinated origins, overlapping root topologies, scope and erasure, temporal bounds, conflict, invalidation, supersession, and emission-head races, alongside five supported controls. Table 5 (Section 5) specifies the per-case oracle, fixture roots, and mediation dispositions.

Table 9: Architectural comparison of objects, enforcement surfaces, and provenance contracts.

Framework	Primary Object	Enforcement Surface	Provenance Contract	Model-Replacement Invariant
Continuity Kernel [7]	State-transition proposal	Accepted-head activation	Proposal / parent hash chain; epistemic support outside contract	State protocol remains invariant across backbones
MemIR [13]	Typed memory atoms	Memory write and projection	Evidence / cue / claim roles; source-root graph not modeled	Not evaluated
MemTX [8]	Belief record & action	Commit gate & action gate	Record dependency lineage; source-root independence not enforced	Multi-backbone support; no fixed-head invariant
ConsistencyGate [20]	Candidate memory fact	Write-time filter	Self-consistency score; does not track upstream source roots	Multi-backbone support; no fixed-head invariant
MiniCheck [21]	Output claim atom	Post-generation gate	Document grounding check; source-root graph not modeled	Not evaluated
This paper	Claim vertex & response unit	Typed resolution & release gate	Exact versioned dependency DAG with policy-relative source roots	Fixed head determines the accepted autobiographical claim set

Table 10: Comparison of temporal validity, correction models, disclosure scoping, and release contracts.

Framework	Temporal / Correction Model	Disclosure Scoping	Assertion Release Contract	Withholding Privacy Safety
Continuity Kernel	Fresh parent head; forward state succession	Outside stated contract	Outside stated contract	Outside stated contract
MemIR	Temporal grounding; correction not modeled	Retrieval / provenance scope	Normalized fact projection; non-bypass not specified	Not specified
MemTX	Record validity & typed cascade repair	Action-focused permissions	Outward assertion release outside contract	Not specified
ConsistencyGate	Temporal validity not modeled	Outside stated contract	Outside stated contract	Not specified
MiniCheck	Temporal validity not modeled	Outside stated contract	Claim-level grounding evaluator	Not specified
This paper	Validity intervals, transitive invalidation, and emission-head checks	Granular principal, purpose, and relationship scope	Decoupled typed content and status release checks	Recipient-safe uniform status rendering

E Sociotechnical Governance, Privacy, and Erasure

E.1 Consent Revocation and Existence Leakage

Proposition 3 (Existence Side Channel under Distinguishable Rendering). *Let $\mathcal{O}(Q, c)$ be a deterministic output class observable to an adversary who knows the rendering policy and can distinguish the classes used for Unknown and Withheld. If those internal branches always map to distinct output classes, observing $\mathcal{O}(Q, c)$ reveals which branch occurred for that query.*

When policy requires those branches to remain indistinguishable, Equation 21 maps the corresponding decisions to $\perp_{unk}$. Equation 26 then gives that common view one render set, closing this branch-distinguishing channel under the proposition’s observation assumptions.

E.2 Relational Privacy and Multi-Party Contexts

A deployment may enforce three relational privacy rules: (1) **Isolation**: claim objects partition private statements from shared context; (2) **Relational Consent**: disclosing p_j ’s statements to p_i requires authorization from p_j ; and (3) **Asymmetric Revocation**: admitted consent changes update later projections without rewriting historical state.

E.3 Erasure-Aware Lineage Design

One possible deployment pattern separates evidence retention from lineage metadata: encrypted evidence may be made inaccessible through key destruction; affected nodes then transition to **Unavailable**; derived claims are recomputed under Equation 22; and non-sensitive lineage metadata may be retained where policy permits. Whether this pattern satisfies GDPR Article 17 or another erasure obligation depends on the jurisdiction, controller obligations, data model, and implementation. The present artifact neither implements cryptographic shredding nor establishes legal compliance.

F Related Work Comparison

Tables 9 and 10 summarize the comparison with selected related systems.