

Where Does Exactly-Once Live?

Model, Harness, and Tool-Contract Effects on Duplicate Side Effects in LLM Agents

Jiapeng Li
Microsoft

jiapengli@microsoft.com

Abstract

When a tool-using agent’s write times out or returns a server error, the action may already have taken effect. Retrying blindly duplicates it—a second charge, a second announcement, a second deployment—while giving up skips required work. We ask where exactly-once behaviour should be enforced: in the model, in the agent harness, or in the tool contract. We introduce LIMBO, a deterministic sandbox of six services with realistic contracts (optional idempotency keys, eventually consistent and missing read paths) and twelve fault modes injected at the service boundary, including late commits, redelivery and partial batches; every episode is graded against a ledger of committed effects. Across 25,930 episodes spanning nine recent models, three production agent harnesses, two contract variants and fifteen recovery conditions, the answer depends on the fault. When an immediate read-back can reveal what happened, the model decides: frontier models instructed to act exactly once almost never duplicate a write whose acknowledgement was lost (0.5%), weaker models often do, and the model explains 53% of the explained variance. When it cannot—the request is still in flight, or the transport delivered it twice—the same frontier models duplicate in 56% and 74% of episodes, and the contract explains 81%. We prove that no verification-only policy is exactly-once under late commits without a bound on in-flight time. Waiting works when such a bound is short and known, but with heavy-tailed in-flight delays even an hour of waiting per episode falls short of offering an idempotency key on every write, which lowers the duplicate rate from 28% to 4% because agents use keys when they exist. The harness barely matters, a guard that attaches keys transfers across harnesses unchanged, and agents reported success in 90% of the episodes in which they had duplicated an effect.

1 Introduction

Language-model agents increasingly act on external systems: they publish posts, charge cards, send email, open tickets and trigger deployments through tool calls. Those calls cross a network, and networks lose messages. When a write times out or returns a server error, the action may or may not have taken effect. A client that retries blindly can charge a customer twice or announce a release twice; a client that gives up can silently skip required work. Distributed systems treat this ambiguity as fundamental—a client cannot distinguish a lost request from a lost acknowledgement—and have converged on remedies that live in the *interface* rather than in the client: idempotency keys, conditional writes and queryable operation status (Helland, 2012; Stripe, 2026; Featonby, 2020).

Agents inherit the ambiguity, but not by default the remedies. Concurrent work has begun to document the consequences. Sun (2026) constructs observation-equivalent counterfactual pairs—identical timeouts with and without a hidden commit—and finds that five open and API models retry blindly about half the time unless they are given a status query or a stable key. IdempotencyBench (Gopnal Swamy, 2026) argues from scripted agents that exactly-once behaviour is a property of the execution substrate rather than of prompting, and Mansoor et al. (2026) show that a verify-then-retry wrapper around a single model lowers duplicate rates. These studies leave the central engineering question open: *where should exactly-once be*

enforced? If frontier models reliably verify before they retry, better models suffice. If agent harnesses—the command-line agents and frameworks that wrap a model—retry transparently or change what the model sees, the harness matters. If some failures cannot be resolved by any amount of verification, the remedy must be in the tool contract.

We study this question with LIMBO, a benchmark in which every recovery decision is graded against a ground-truth ledger of committed side effects, and with a factorial experiment that varies the model, the harness, the tool contract and the recovery policy while holding the task, the world and the fault fixed. LIMBO injects twelve fault modes at the service boundary, including three that prior agent benchmarks omit: *late commits*, where a timed-out request is still in flight and lands after the client has moved on (with a fixed or a heavy-tailed delay); *redelivery*, where an at-least-once transport delivers a request twice; and *partial batches*. Its tools carry realistic contracts—optional idempotency keys, eventually consistent read paths with documented lag, endpoints with no read-back at all, naturally idempotent operations—so that the same fault can be recoverable, recoverable only with care, or unrecoverable depending on the contract. Because the sandbox is exposed through the Model Context Protocol, the identical environment runs under a minimal function-calling scaffold and under production agent harnesses.

Our study covers nine recent models from five providers in a minimal scaffold and three production harnesses (GitHub Copilot CLI, Hermes and Codex CLI), for a total of 25,930 graded episodes. The answer to our question depends on the kind of failure, and we find:

- **When a read-back can resolve the failure, the model decides.** Frontier models instructed to act exactly once almost never duplicate a write whose acknowledgement was lost (0.5% of episodes), because they verify before retrying; weaker models duplicate in 18%, and without the explicit instruction the weaker and faster models duplicate markedly more while the strongest are unaffected. On these faults the model explains 53% of the explained variance in duplicates.
- **When it cannot, only the contract does.** When the timed-out request is still in flight, the same verification finds nothing and the retry creates a duplicate in 56% of frontier episodes; under redelivery, in 74%. Here the contract explains 81% of the variance. We prove that no verification-only policy can be exactly-once under late commits without a known bound on in-flight time, and measure the price of waiting for such a bound: waiting works when the bound is short and known, but with heavy-tailed in-flight delays even an hour of waiting per episode falls short of what keys achieve with no added latency. Offering keys is enough: under a counterfactual contract in which every write accepts one, agents attach them in 98% of episodes and the duplicate rate falls from 28% to 4% with no other change; every remaining duplicate involves an agent that sent the first attempt without a key or changed the key when it retried.
- **Harnesses matter little; client-side middleware has a ceiling.** Three production harnesses and a minimal scaffold running the same model behave almost identically. Transparent client-side retries cut exactly-once success from 72% to 50%. In the native contract even an outcome oracle that sees in-flight requests reaches only 87%, because redelivery on key-less writes is beyond any client-side policy; with keys on every write and a guard that attaches them, exactly-once success reaches 99%. The same holds in every harness: with keys everywhere, the duplicate rate of a shared model is at most 0% even without the guard, and exactly-once success is at least 100% with it.
- **Agents do not know when they have duplicated.** In 90% of episodes that produced a duplicate, the agent reported the task as completed.

We release the benchmark, the harness adapters and every episode trace.¹

2 Related work

Ambiguous tool outcomes. Three concurrent efforts study agents facing tool calls whose effect is unknown. Sun (2026) introduces observation-equivalent counterfactual pairs over 81 software engineering templates and compares plain prompting, prompt advice, a status query and a stable idempotency key across five

¹Code and data will be made available upon publication.

model endpoints; eventual consistency and partial commits are explicitly out of scope and no mitigation is proposed. IdempotencyBench (Gopnal Swamy, 2026) measures idempotency violations with a hidden effect ledger under several retry regimes, reporting results for scripted agents. Mansoor et al. (2026) evaluate one model on two tasks under timeouts, delayed visibility and partial success, and show that a wrapper combining post-condition verification, idempotency keys and verify-before-retry reduces duplicates; the wrapper resolves ambiguity without informing the model. We adopt the counterfactual pairing of Sun (2026) as a design principle, reproduce the verify-before-retry wrapper as a baseline, and add what these studies leave open: late commits and redelivery, production harnesses, many recent models, a contract manipulation, and a belief-visible guard that also escalates.

Tool failures and agent robustness. Hell or High Water (Wang et al., 2025) evaluates recovery when tools become unavailable and alternatives exist; ToolMaze (Zhu et al., 2026) studies replanning under explicit and implicit perturbations; ReliabilityBench (Gupta, 2026) adds timeouts, rate limits, partial responses and schema drift; WAREX (Kara et al., 2026) injects web and network faults; AgentCheck (Mazumder & Lia, 2026) and ToolMisuseBench (Sigdel & Baral, 2026) provide deterministic, replayable fault injection; AgentChaos (Tan et al., 2026) injects faults into LLM API responses of multi-agent systems. PALADIN (Vuddanti et al., 2025) trains recovery from tool failures. These works measure task completion under observable failures; none tracks whether a failed write actually took effect. Diagnostic studies of long-horizon failures (Wang et al., 2026b; Zhu et al., 2025a) and self-correction methods (Shinn et al., 2023; Gou et al., 2024) likewise score recovery by task success, so a duplicate costs nothing.

Transactions, rollback and runtime enforcement. GoEX (Patil et al., 2024) argues for undo and damage confinement in LLM runtimes; SagaLLM (Chang & Geng, 2025), Atomix (Mohammadi et al., 2026), Cordon (Chen et al., 2026), DART (Yang et al., 2026) and Agentic Transaction (Sun et al., 2026) provide transactional or compensating execution; AgentRewind (Zhuang et al., 2026) checkpoints and rewinds agent context and workspace but treats external calls as irreversible and replays them from its log. AgentSpec (Wang et al., 2026a) enforces user-specified rules at runtime. These are mechanisms; we measure how often, and where, a mechanism is needed, and how much a minimal contract-level mechanism buys across models and harnesses.

Stateful benchmarks and reliability measurement. τ -bench (Yao et al., 2025), τ^2 -bench (Barres et al., 2025), AppWorld (Trivedi et al., 2024), ToolSandbox (Lu et al., 2024), ToolEmu (Ruan et al., 2023), AgentDojo (Debenedetti et al., 2024) and MCPMark (Wu et al., 2025) grade agents against world state, which our grader extends with an effect ledger. Rabanser et al. (2026), HAL (Kapoor et al., 2025a) and Kapoor et al. (2025b) argue for measuring reliability and cost rather than accuracy alone, and Zhu et al. (2025b) give a checklist for rigorous agentic benchmarks that informed our validity checks. Tang & Zhan (2026) allocate a verification budget across agent actions by downstream harm; our setting concerns verification after a failure rather than before an action. A source-code study of eleven coding-agent harnesses (Barbaste et al., 2026) documents their architectures; our harness experiments measure the resulting behaviour under identical faults.

3 Recovery under uncertain side effects

Setting. An agent issues a write request r to a service. The service may execute r before answering, after the client has stopped waiting, partially, more than once, or not at all. We call this hidden fact the *outcome state* $s \in \{\text{absent}, \text{committed}, \text{late}, \text{partial}, \text{duplicated}\}$. The agent observes only a response o (an acknowledgement, a timeout, or an error such as HTTP 500) and whatever it later reads back through the service’s query tools. Reads may be eventually consistent: a read at time t reflects only effects committed before $t - \lambda$ for a lag $\lambda \geq 0$, and some services offer no read-back path at all.

Definition 1 (Observation equivalence). *Two outcome states s, s' are observation-equivalent for a response o if $P(o \mid s) = P(o \mid s')$ and every read issued before the effect of r becomes visible returns the same result under s and s' .*




Figure 1: Why verification is not enough. W_{absent} and W_{late} produce identical observations—a timeout and then an empty read-back—until the in-flight original commits, so verify-then-retry re-issues in both and duplicates in W_{late} (Proposition 1). If the re-issue reuses the original’s idempotency key, the late original becomes a no-op in every world (Proposition 2).

A timeout is observation-equivalent under **absent**, **committed** (response lost) and **late** (request still in flight). A recovery policy therefore cannot condition on the outcome state directly; it can only gather evidence (by reading, waiting, or asking a human) or choose actions that are safe under every equivalent state.

Outcomes. For a task with required effects $\{e_j\}$, let x_j be the number of committed effects matching e_j over the whole episode (including effects later undone) and ℓ_j the number still standing at the end. The task succeeds (TS) if every ℓ_j reaches its required count and no unrelated record was damaged; it is *exactly-once successful* (EOS) if, in addition, no x_j exceeds its required count. A duplicate side effect is any excess execution, whether or not the agent later compensated for it: a refunded double charge was still charged.

Why verification is not enough. The natural recovery strategy—read back, and re-issue only if the effect is absent—is exactly-once when the effect of r is visible by the time the agent reads. It fails when the request is still in flight.

Proposition 1 (No verification-only exactly-once under late commits). *Consider a non-idempotent write without key support that times out at time t_0 , and the two worlds W_{absent} , in which r never executes, and W_{late}^δ , in which r executes at $t_0 + \delta$. Let π be any recovery policy whose actions depend only on observations. If π completes the task in W_{absent} with probability one within time T , then for every $\delta > T + \lambda$, π produces a duplicate in W_{late}^δ with probability one.*

Proof. In W_{absent} the task can only be completed by re-issuing r , which π therefore does at some time $t_r \leq t_0 + T$. Until $t_0 + \delta + \lambda$ every read returns the same result in both worlds, so for $\delta > T + \lambda$ the observation histories of the two worlds coincide up to t_r and π re-issues r in W_{late}^δ as well. The in-flight request then commits at $t_0 + \delta$ in addition to the re-issued one. $\square$

A known bound Δ on in-flight processing restores a verification-only solution—wait until $t_0 + \Delta + \lambda$, then read—at a latency cost of $\Delta + \lambda$ per ambiguous write; few APIs document such a bound, none of our tools do, and if the true delay has a heavy tail, any finite assumed bound leaves the mass beyond it exposed. Section 6.3 measures this trade-off. Idempotency keys remove the dilemma altogether.

Proposition 2 (Keys suffice). *If a service executes each idempotency key at most once and returns the original outcome for repeats, and a batch interrupted under a key resumes where it stopped, then the policy “re-issue with the same key until acknowledged” is exactly-once in all five outcome states.*

Proof. Every execution of r , including the in-flight original and any redelivered copy, carries the same key, so at most one execution has an effect; re-issuing until an acknowledgement arrives guarantees at least one. $\square$

Propositions 1 and 2 make the question of this paper concrete and split it in two. For faults that an immediate read-back resolves—a lost acknowledgement, a misleading error, a partial batch—careful verification by the model suffices, so exactly-once depends on the model and on what the harness lets it see. For faults that

Table 1: Services and tool contracts. “Key” marks writes that accept an idempotency key in the native contract; lags are documented in the tool descriptions.

Service	Writes (idempotency)	Read-back path (consistency)
social	publish (key on mastodon only)	list_posts : weibo lags 180 s; none for x
billing	create_charge (key), refund (naturally idempotent)	list_charges (strong)
tickets	create , add_comment (non-idempotent), update_status (conditional)	list_recent (strong), search (lags 120 s)
mail	send (non-idempotent, irreversible)	search_sent (lags 120 s)
data	insert , insert_many (non-atomic), upsert (idempotent)	query (strong)
deploy	trigger (non-idempotent)	list_runs , get_run (strong)

a read-back cannot resolve—a request still in flight, or a request delivered twice—model-side reasoning can at best decline to act (escalate, or report the operation as uncertain); it cannot complete the task exactly once on a key-less write. There, exactly-once depends on the contract (does the write accept a key?) and on whether someone attaches the key. The rest of the paper measures each factor in each regime.

4 The Limbo benchmark

LIMBO is a deterministic, resettable sandbox of six simulated services, twelve parameterized task templates, a fault injector that operates at the service boundary, and a grader that reads only the ground-truth effect ledger. Table 1 summarizes the services.

Services and contracts. Semantics follow well-known production conventions. Billing keys follow Stripe: a repeated key with the same parameters replays the original charge and a repeated key with different parameters is rejected (Stripe, 2026). Social posts, tickets, emails, database rows and deployment runs are created anew by every successful call unless a key is honored. Several read paths are eventually consistent with a documented lag (a search index, a mail Sent folder, one social platform’s listing), and one platform has no listing endpoint, so its writes cannot be verified. Search endpoints behave like search engines: a record matches when every query term occurs in its fields. Three general tools complete the interface: **wait** advances a virtual clock; **escalate_to_human** reaches a simulated on-call operator who inspects the ground truth of every write that returned an error and answers after 15 simulated minutes (or, in one condition, never answers); and **finish** records a structured final report with a status and a list of operations whose outcome the agent considers uncertain. Every tool additionally carries a machine-readable contract in the style of MCP tool annotations (Model Context Protocol, 2025): read-only, idempotent and destructive hints, the fields that identify a write’s intent, a read-back function, the documented visibility lag and a compensating action. The agent never sees the contract object; only the harness-level recovery policies in §5 may read it.

Tasks. Twelve templates describe realistic operational workflows: announcing a release on two platforms, opening a ticket with the post links and emailing a summary; billing two or three customers for an invoice and notifying finance; opening an incident and paging on-call; recording a batch of database migrations; rolling out a service to staging and then production; refunding a duplicate charge while keeping the original; enabling feature flags and announcing them; notifying three customers individually; cross-posting an event; upgrading a subscription; resolving an incident ticket; and a long-horizon hotfix rollout with about ten writes across five services. Each template is instantiated with seeded parameters (products, versions, customers, amounts) and states its targets precisely (for example, *exactly one* post containing the release version on each named platform). Each instruction ends with one sentence asking for every action to happen exactly once; E6 removes it to measure how much this cue drives caution. Each template also declares its *focal writes*—the calls a fault may attach to—annotated with idempotency class (non-idempotent, key optional, naturally idempotent, idempotent), verification class (strong, eventual, none) and reversibility. Across templates there are 34–35 focal writes per instance.

Fault injection. A fault is attached to the n -th call that matches a focal write’s tool and argument filter, so the same world and the same fault apply to every model, harness and condition. Table 2 lists the twelve modes. The observable response of each ambiguous mode is byte-identical across its hidden outcome states,

Table 2: Fault modes. Rows in the same group are observation-equivalent at the time of the response.

Mode	Agent observes	Hidden outcome	Exactly-once recovery
<code>timeout_pre</code>	timeout	not executed	verify, then re-issue
<code>timeout_post</code>	timeout	executed, response lost	verify, then skip
<code>timeout_late</code>	timeout	executed 90 s later (still in flight)	same key, or wait and escalate
<code>timeout_late_tail</code>	timeout	executed 40 s–2 h later (heavy tail)	same key
<code>http500_pre</code>	HTTP 500	not executed	verify, then re-issue
<code>http500_post</code>	HTTP 500	executed	verify, then skip
<code>partial_timeout</code>	timeout on a batch	first half executed	verify, re-issue only missing rows
<code>duplicate_delivery</code>	success	executed twice	send a key proactively
<code>http503_transient</code>	503, retry later	not executed	retry
<code>rate_limit</code>	429 with retry-after	not executed	wait, then retry
<code>outage</code>	persistent 503	not executed	stop and report
<code>schema_drift</code>	400, field renamed	not executed	repair arguments
<code>none</code>	success	executed	–

and the virtual clock advances identically, so a pair such as `timeout_pre/timeout_post` differs only in whether the service committed. A request that the service would reject anyway receives its real error and leaves the trigger armed for the next matching call; this keeps the pairs symmetric even when an agent’s first attempt is malformed. Late commits land 90 simulated seconds after the request was sent (`timeout_late`) or after a delay drawn log-uniformly between 40 s and 2 h (`timeout_late_tail`, median about 9 minutes), whatever the agent does in the meantime; the delay is drawn deterministically per world, so every model and policy faces the same one. At the end of an episode, requests still in flight complete, as they would in reality.

Grading. The grader reads the ledger of committed effects and the final world state; no model is used as a judge. It reports task success (TS), exactly-once success (EOS), the number of duplicate executions (including compensated ones), residual duplicates, collateral damage to records the task did not target (for example refunding the original charge), unrequested writes, and *overclaim*: finishing with status `completed` although the task failed or a duplicate remains. From the agent’s own tool calls—not the harness’s—a deterministic coder classifies the recovery behaviour that follows the focal fault as blind retry, retry with the same key, retry with a new key, verify then retry, verify then skip, escalate, stop without checking, or move on.

Validity checks. Thirty-one unit tests with scripted agents check the fault semantics exactly: observation equivalence of each pair, symmetric handling of invalid requests, stale reads before and fresh reads after the documented lag, partial batches, key replay, conflict and batch resumption, late commits defeating verification but not keys, the heavy-tailed delay distribution, the waiting thresholds of the wait-then-verify policies, redelivery, escalation reporting ground truth, and the grader’s detection of collateral damage and overclaiming. Scripted reference policies behave as the propositions predict: a blind-retry agent duplicates only in committed worlds, a verify-first agent is exactly-once except under late commits, and a same-key agent is exactly-once everywhere. In the empirical data, agents cannot distinguish the observation-equivalent worlds (§6), confirming that no information about the hidden outcome leaks through the response.

Harness-agnostic deployment. The sandbox runs as a local HTTP service owned by the experiment runner. A stdio MCP server forwards only `tools/list` and `tools/call` to it, so any MCP-capable agent can be evaluated unchanged; the fault plan and the ground truth never leave the runner. Each harness episode runs in an empty working directory with a throwaway harness home, so memory and session state cannot carry over between episodes.

5 Experimental setup

Models. We evaluate nine models reachable through one gateway, spanning five providers and several capability tiers: `gpt-6-astra`, `gpt-6-sol`, `gpt-5.6-sol`, `gpt-5.4-mini` and `gpt-4.1` (OpenAI), `claude-opus-5.5` (Anthropic), `gemini-3.8-flash` (Google), `grok-4.7` (xAI) and `mai-code-1.1-flash` (Microsoft). All models are used with provider-default sampling and reasoning settings. Model identifiers are those exposed by the gateway at the time of the study (September 2026).

Minimal scaffold. The scaffold is a standard function-calling loop (Yao et al., 2023): a fixed system prompt (Appendix A.2), the task as the first user message, the tools of the task’s services plus the three general tools, and up to 40 tool calls. Parallel tool calls are executed in order. If a turn contains no tool call, the scaffold asks the model once more to continue or finish; after three such turns the episode ends. The same conversation is rendered to the chat-completions or the responses protocol depending on the model, with reasoning items replayed where the protocol supports it.

Production harnesses. We run three agent command-line interfaces non-interactively in an empty temporary directory with LIMBO as their only MCP server: GitHub Copilot CLI 1.0.86, Hermes Agent 0.20.6 and OpenAI Codex CLI 0.139. Each harness keeps its own system prompt, tool-call protocol, context management and retry logic; we prepend the same operational preamble to the task (Appendix A.2), restrict the harness to the sandbox’s tools where it allows this, and give each episode a fresh harness home directory. All harnesses reach their models through the same gateway as the scaffold, which lets us cross harnesses with models.

Contract variants. In the *native* contract (Table 1) only billing and one social platform accept idempotency keys. The counterfactual *keys-everywhere* contract extends Stripe-style key semantics to every non-idempotent write, including resumable batches, late commits and redelivered requests, and documents the new parameter in each tool description. Nothing else changes.

Recovery conditions. Two conditions change only the prompt: *aware* adds five sentences of reliability guidance to the system prompt, and *reflect* inserts a generic reflection request after any turn with a tool error (Shinn et al., 2023). The remaining conditions wrap tool execution, as a framework or middleware would, and never see ground truth: *sdk-retry* transparently retries timeouts, 5xx and 429 responses up to three times with exponential backoff, mimicking common client defaults; *rules* retries only 429 and 503 responses; *vbr* (verify-before-retry, after Mansoor et al. (2026)) intercepts a write whose identical predecessor returned an ambiguous error and runs its declared read-back first, sending the write only if the effect is absent; *wait* Δ (E5) additionally assumes an in-flight bound of Δ seconds and defers that read-back until Δ seconds, plus the documented visibility lag, after the original request was sent; and *guard* adds four contract-driven components to *vbr*—automatic idempotency keys where the contract accepts them, re-verification after the documented visibility lag, blocking an unverifiable non-idempotent repeat with a request to escalate, and an annotation telling the model that an outcome is unknown. Two reference conditions may read ground truth. The *state oracle* is the guard verifying against the current ground-truth state; it cannot see a request that is still in flight, so it is *not* an upper bound under late commits. The *outcome oracle* also sees in-flight requests and suppresses any re-issue of a write that will commit; it is the upper bound for any client-side policy, since only redelivery on writes without key support remains beyond its reach. In harness runs the wrappers execute inside the MCP server, so they apply to any harness without modification.

Experiments. E1 crosses the nine models with the eleven fault modes of the original design (all but the heavy-tailed late commit) and every focal write of two instances per template in the minimal scaffold, plus a fault-free episode per instance (7,484 episodes with the supplement below), and adds instances for the two sparsest contract classes (partial batches and the unverifiable endpoint); `gpt-4.1` is served from a separate low-rate quota and was run in its own low-concurrency lane with the identical design. E2 crosses four models (`claude-opus-5.5`, `gpt-6-sol`, `gemini-3.8-flash`, `mai-code-1.1-flash`) with nine conditions (vanilla, aware, reflect, sdk-retry, rules, vbr, guard and the two oracles) and the eight fault modes most relevant to recovery, and repeats vanilla and guard under the keys-everywhere contract (E2k; 9,548 episodes in total). E3

runs Copilot CLI and Hermes with `claude-opus-5.5`, `gpt-5.6-sol` and `gemini-3.8-flash`, and Codex CLI (which speaks only the responses protocol) with `gpt-5.6-sol` and `gpt-6-sol`, natively and with the guard in the MCP server, on the core fault modes and two focal writes per template; the minimal scaffold is run on the identical design as the reference harness. E3k repeats the harness design under the keys-everywhere contract with one model shared by all four harnesses (`gpt-5.6-sol`; 3,968 episodes for E3 and E3k). E4 probes robustness on two focal writes per template, compared with matched E2 episodes: tool descriptions without consistency statements or with explicit non-idempotency warnings, three paraphrases of the system prompt, an operator who never answers, and guard ablations (1,928 episodes). E5 runs the wait-then-verify family ($\Delta \in \{0, 60, 120, 300\}$ s under the fixed delay; $\{0, 60, 300, 900, 3600\}$ s under the heavy-tailed delay) against vanilla, the guard, the outcome oracle and the keys-everywhere contract on the four E2 models and all focal writes of instance 0 (1,904 episodes). E6 removes the closing “exactly once” sentence from every instruction for six models on instance 0 and pairs each episode with its E1 counterpart (1,098 episodes).

Statistics. The world seed depends only on (template, instance, focal write, fault mode), so all models, harnesses, conditions and instruction variants face identical worlds, and comparisons are paired. Episodes that share a task template are correlated, and there are only twelve templates; a cluster bootstrap over so few clusters is anti-conservative and degenerate for cells with no events. Per-cell 95% intervals are therefore Wilson intervals on a Kish effective sample size $n/(1 + (\bar{m} - 1)\rho)$, where $\bar{m}$ is the mean number of episodes per template in the cell and ρ the within-cell intra-class correlation by template, estimated once per experiment. Regression analyses use a Bayesian binomial mixed model with random intercepts for template and template $\times$ instance (as preregistered for H1, fitted by variational Bayes) and, as a frequentist check, generalized estimating equations clustered by template with bias-reduced (Mancl-DeRouen) standard errors and a t reference with $G - 1$ degrees of freedom. Paired binary outcomes use exact McNemar tests with Holm correction within each family; H4 uses the preregistered two one-sided tests (TOST) with a ± 5 pp margin; observation equivalence is assessed by comparing first-action agreement across matched worlds with agreement across independent runs of the same world, and by a within-pair permutation test of the total variation distance. The Shapley decomposition of McFadden’s pseudo- R^2 is reported pooled, as preregistered, and separately for faults that an immediate read-back resolves and faults it cannot. Episodes whose focal write was never attempted are excluded from fault-conditional analyses and reported as trigger rates; episodes that end in a gateway error after eight attempts are excluded and re-run. Hypotheses, metrics and analyses were preregistered before E1; E2k, E3k, E5, E6, the outcome oracle, the stratified decomposition and the revised interval method were added afterwards and are exploratory, and every deviation from the preregistration is listed in the released materials.

6 Results

6.1 What frontier models get right, and where they fail (RQ1)

Figure 2 and Table 3 summarize E1. The six frontier models almost never duplicate a write whose acknowledgement was lost: under `timeout_post` their duplicate rate is 0.5%, and under a misleading HTTP 500 it is 8%. The mechanism is visible in their behaviour (Figure 6): after a timeout they overwhelmingly read back before acting. The two weaker models (`mai-code-1.1-flash` and `gpt-5.4-mini`) fail even this textbook case, duplicating in 18% of lost-acknowledgement and 47% of misleading-500 episodes. Pooled statistics exclude `gpt-4.1`, which is served from a separate low quota and completed only 87% of its E1 design; it is shown per model with this caveat.

The same frontier models duplicate in 56% of late-commit and 74% of redelivery episodes. Late commits defeat exactly the behaviour that protects them elsewhere, and the traces show how. Of the 297 late-commit episodes in which the agent re-issued the write within a minute of the timeout, 69% produced a duplicate: the read-back happened while the original was still in flight. Waiting did not help as agents implemented it: all 123 agents that waited at least 60 s before re-issuing duplicated (100%), and 99% of them were on eventually consistent read paths, where they waited the documented visibility lag *measured from the error* rather than from the unknown moment at which the in-flight request would commit. Agents that never




Figure 2: Duplicate side-effect rate (%) by model and fault mode (minimal scaffold, vanilla prompt, E1). Columns are grouped as in Table 2; rows are ordered roughly by capability. The partial-batch column includes the EIs supplement (ten instances of the batch template per model).

Table 3: E1 results by model with 95% intervals (Wilson on Kish effective sample sizes). Duplicate rates are conditional on the fault having fired; TS (not executed) is task success in the `_pre` worlds; EOS is over all faulted episodes. `gpt-4.1` covers only 87% of the design (separate low quota).

Model	Duplicate rate (%) under committed faults			TS (%) when not executed	EOS (%) all faults
	lost ack / 500	partial	late commit		
gpt-6-astra	0 [0,9]		25 [13,42]	74 [57,86]	100 [91,100]
gpt-6-sol	0 [0,9]		46 [30,63]	74 [57,86]	99 [89,100]
gpt-5.6-sol	1 [0,10]		65 [48,79]	74 [57,86]	98 [88,100]
claude-opus-5.5	0 [0,9]		68 [51,82]	74 [57,86]	100 [91,100]
gemini-3.8-flash	14 [6,29]		59 [42,75]	74 [57,86]	100 [91,100]
grok-4.7	11 [4,24]		72 [55,85]	75 [58,87]	100 [91,100]
mai-code-1.1-flash	43 [28,59]		70 [52,83]	75 [58,87]	93 [80,98]
gpt-5.4-mini	21 [11,37]		71 [54,84]	75 [58,87]	100 [91,100]
gpt-4.1	83 [68,92]		84 [66,93]	85 [68,94]	98 [87,100]

re-issued the write were exactly-once in all but 1% of 84 episodes. Redelivery never produces an error at all, so only an idempotency key sent in advance can prevent it.

The benchmark does not reward caution for its own sake. When the request demonstrably never executed (`timeout_pre`, `http500_pre`), agents duplicate in only 0.6% of episodes and complete the task in 99%; explicit, non-ambiguous faults (503, 429, outage, schema drift) cause duplicates in 0.4% of episodes, and fault-free episodes in 0.0%. Agents also cannot tell the observation-equivalent worlds apart, as intended. On matched pairs of worlds that differ only in whether the timed-out write committed, the agent’s first action after the timeout agreed in 87% of pairs, no less often than two independent runs of the *same* world agree (82%; $n = 136$ and 272 pairs), and the total variation distance between the two worlds’ first-action distributions (0.07) equals that expected under exchangeable labels (0.07; paired permutation $p \geq 0.15$ for every model). Outcomes are stable across independent runs: on the 868 episodes run in both E1 and E2 with identical worlds, duplicate outcomes agree in 96% of cases (Cohen’s $\kappa = 0.90$).

6.2 Where exactly-once lives depends on the fault (RQ2)

The faults in Table 2 fall into two classes that call for different remedies. For a lost acknowledgement, a misleading 500 or a partial batch, an immediate read-back reveals what happened, so a careful agent can recover without help from the service. For a late commit or a redelivered request it cannot: the read-back

Table 4: Shapley decomposition of the explained variance in duplicates. “Pooled” is the preregistered analysis over committed faults and redelivery; the other columns split it into faults that an immediate read-back resolves (lost acknowledgement, misleading 500, partial batch) and faults it cannot resolve (late commit, redelivery). Top: E1 (minimal scaffold, eight models). Bottom: E3 (four harnesses, native condition).

Factor	pooled (prereg.)	read-back resolves	read-back cannot
contract share (%)	36	30	81
mode share (%)	56	18	10
model share (%)	9	53	8
total pseudo- R^2	0.66	0.51	0.72
duplicate rate (%)	38	11	66
episodes	3,312	1,696	1,616
Factor	pooled (prereg.)	read-back resolves	read-back cannot
contract share (%)	35	42	88
mode share (%)	64	17	9
model share (%)	1	39	2
harness share (%)	0	3	0
total pseudo- R^2	0.77	0.54	0.76
duplicate rate (%)	35	4	67
episodes	1,164	588	576

either precedes the effect or is never prompted by an error. Table 4 decomposes the explained variance in duplicates (Shapley shares of McFadden’s pseudo- R^2) separately for the two classes.

When a read-back resolves the fault, the model decides. On resolvable faults the model accounts for 53% of the explained variance and the contract for 30% (E1). This is the regime in which frontier models are nearly exactly-once and weaker models are not. In E3, whose models are all frontier-class, the model and the contract contribute similarly on these faults (39% and 42%), because the frontier models differ little from one another here.

When it does not, only the contract does. On late commits and redelivery the contract accounts for 81% and the model for 8%. In the harness experiment the harness adds 3% on resolvable and 0% on unresolvable faults. The preregistered pooled analysis mixes the two regimes and therefore attributes 36% to the contract, 56% to the fault stage and only 9% to the model (H2 supported as preregistered, but the stratified analysis shows that this is a statement about the unresolvable faults; the preregistered harness share of at least 10% is not supported).

Contract class. Figure 3 breaks duplicates down by the contract of the focal write. Among frontier models, a write that accepts an idempotency key is duplicated in 1.5% of committed-fault episodes and an idempotent write in 0.0%; frontier models attach a key to a keyable write in 98% of these episodes without being told to, the weaker models in 85%. Non-idempotent writes are duplicated in 28% of episodes when they can be read back from a strongly consistent endpoint and 31% when the read path is eventually consistent. In the preregistered mixed-effects logistic regression (random intercepts for template and instance, fixed effects for model and fault), a non-idempotent write with an eventual or missing read path has 6.4 times the odds of being duplicated as a keyable, idempotent or strongly verifiable one (95% credible interval 5.2–7.9); a GEE with bias-reduced cluster-robust errors gives 2.9 (1.4–6.0, $p = 0.007$). H1 is supported.

A read path that cannot see lagging effects is worse than none. Writes with no read-back path at all are duplicated less often (13%) than writes that can be read back. Without a way to check, frontier agents escalated to the operator in 87% of committed-fault episodes, and the operator resolves the uncertainty; when a read path exists they verified instead (96% strong, 89% eventual) and trusted a read that could not see a lagging or in-flight effect. Restricting to lost-acknowledgement faults in which the agent read back before




Figure 3: Duplicate rate under committed faults (lost acknowledgement, misleading 500, late commit, partial batch) by tool-contract class and model (E1 and E1s, with 95% intervals).

Table 5: Duplicate rate (%) by contract variant and condition under committed faults and redelivery (E2, E2k; four models pooled).

Contract	Condition	timeout (committed)	500 (committed)	timeout (late commit)	partial batch	redelivery
native	vanilla	9	21	61	25	74
native	guard	0	0	68	0	74
keys-everywhere	vanilla	4	6	9	0	7
keys-everywhere	guard	0	0	7	0	0

re-issuing, the duplicate rate was 13.4% on eventually consistent and 0.8% on strongly consistent read paths ($n = 947$; GEE $p = 0.005$; H3 supported).

Changing the contract. The keys-everywhere contract isolates the causal effect of the contract (Table 5). With keys available on every write but no guard, vanilla agents attach a key in 98% of faulted episodes, and the late-commit duplicate rate falls from 61% to 9% and the redelivery rate from 74% to 7%. With the guard attaching keys automatically, late-commit duplicates fall from 68% to 7%, redelivery duplicates from 74% to 0%, and exactly-once success reaches 99%.

6.3 Waiting is not a substitute for keys (E5)

Proposition 1 allows one client-side escape from late commits: if the in-flight delay has a known bound Δ , waiting Δ (plus the read-path lag) before verifying is exactly-once. E5 measures what that costs. A family of harness policies *wait* Δ intercepts an identical re-issue of an unknown-outcome write, waits until Δ seconds (plus the documented lag) after the original request was sent, and verifies before letting the write through. We run them under the fixed 90 s delay of E1 and under a heavy-tailed delay drawn log-uniformly between 40 s and 2 h (median 6 min), and compare them with the guard, with the outcome oracle that sees in-flight requests, and with the keys-everywhere contract (Table 6, Figure 4).

With a fixed delay, waiting works once the assumed bound exceeds the true one: *wait 60 s* reaches 43% exactly-once success and *wait 120 s* 99%, at 3.8 simulated minutes per episode against 2.9 for vanilla—more than the keys-everywhere contract achieves as agents use it (93% with the guard). A short, known bound therefore makes waiting a good remedy. With a heavy-tailed delay the same policies trade latency for success without closing the gap: *wait 5 min* reaches 67% and *wait 1 h* 84% at 49.9 minutes per episode. A one-hour bound covers the in-flight delay in only 82% of worlds, and in the remaining worlds every re-issue duplicates; only episodes in which the agent never re-issued escape. Under the same delays the keys-everywhere contract with the guard reaches 94% at 1.5 minutes. The outcome oracle, which waits exactly as long as necessary

Table 6: Waiting versus keys under late commits (E5; four models, all focal writes of instance 0). “time” is the mean agent-visible episode time in simulated minutes; $P(\delta \leq \Delta)$ is the fraction of tail-delay worlds whose in-flight delay does not exceed the assumed bound. Fixed-delay rows for vanilla, guard, both oracles and the keys contract are the matched E2/E2k episodes.

Condition	fixed delay (90 s)			heavy-tailed delay (40 s–2 h)			
	EOS (%)	dup. (%)	time (min)	EOS (%)	dup. (%)	time (min)	$P(\delta \leq \Delta)$
vanilla	39	61	2.9	34	66	2.8	–
wait 0 s	38	62	2.8	36	64	2.8	0
wait 60 s	43	57	3.2	42	58	2.9	12
wait 120 s	99	1	3.8	–	–	–	–
wait 300 s	99	1	6.2	67	33	6.1	50
wait 900 s	–	–	–	71	29	14.3	62
wait 3600 s	–	–	–	84	16	49.9	82
guard	32	68	2.2	32	68	2.3	–
state oracle	54	46	1.9	–	–	–	–
outcome oracle	100	0	2.3	99	1	19.5	–
keys: vanilla	91	9	0.9	92	8	0.9	–
keys: guard	93	7	1.5	94	6	1.5	–




Figure 4: Exactly-once success against mean agent-visible episode time for wait-then-verify policies with increasing assumed in-flight bounds, compared with keys (E5). Left: fixed 90 s delay. Right: heavy-tailed delay.

because it knows when the in-flight request will commit, reaches 99% at 19.5 minutes: the best that any waiting policy can do.

Keys fall short of the oracle only because agents do not always reuse them. Among the 432 late-commit re-issues under the keys-everywhere contract (E2k, E5k) in which the agent reused the original key, 0% produced a duplicate. The duplicates came from the 56 re-issues whose first attempt carried no key (68% duplicated) and the 2 in which the agent generated a new key for the retry, for example by appending `-retry1` (100%): in both cases the in-flight original and the retry carry different keys. A harness that pins one key per intent, replacing any key the model supplies on a re-issue, would remove this failure mode. This corroborates, with real models, the observation of Gopnal Swamy (2026) from scripted agents that idempotency keys protect only when they are stable across retries.

6.4 Agents do not know what happened (RQ3)

Of the 1,279 E1 episodes that produced at least one duplicate, the agent finished with status `completed` in 90% (95% interval 78–96%), and in 80% it also listed no operation as uncertain (H5 supported). A human reading the final report would have had no reason to check. Agents were also not more careful with irreversible writes: after an ambiguous fault they re-issued an irreversible write (an email or a comment)

Table 7: Recovery conditions (E2). EOS, duplicate rate and TS are over faulted episodes; the bottom rows report mean cost per episode over all four models. The state oracle verifies against the current ground truth and cannot see requests still in flight; the outcome oracle can, and is the upper bound for any client-side policy.

Model	vanilla	aware	reflect	sdk-retry	rules	vbr	guard	state oracle	outcome oracle
<i>EOS (%)</i>									
gpt-6-sol	79	78	78	51	80	78	76	82	88
claude-opus-5.5	77	77	77	51	77	77	77	82	88
gemini-3.8-flash	74	77	71	51	74	73	76	81	88
mai-code-1.1-flash	59	67	58	48	59	62	76	74	83
<i>Duplicate rate (%)</i>									
gpt-6-sol	20	22	22	49	20	22	24	18	12
claude-opus-5.5	23	23	23	49	23	23	23	18	12
gemini-3.8-flash	26	23	29	49	26	27	24	19	12
mai-code-1.1-flash	40	32	42	52	40	35	23	24	15
<i>TS (%)</i>									
gpt-6-sol	100	100	100	100	100	100	100	100	100
claude-opus-5.5	100	100	100	100	100	100	100	100	100
gemini-3.8-flash	100	100	100	100	100	100	100	100	100
mai-code-1.1-flash	100	100	99	100	99	98	99	99	98
mean tool calls	8.7	9.3	8.4	6.6	8.3	8.7	8.6	8.7	8.7
mean tokens (k)	15.5	18.7	15.1	10.8	14.7	15.7	15.8	16.0	14.5
mean sim. time (min)	1.7	1.6	1.5	0.7	1.5	1.6	1.6	1.4	1.4
mean human (min)	0.6	0.3	0.4	0.0	0.4	0.5	0.4	0.2	0.2

without checking in 9.6% of episodes and a reversible one in 6.3% (difference +3.4 pp, 90% interval -1.6 to +8.5 pp, $p = 0.25$). The preregistered TOST is inconclusive: equivalence within ± 5 pp cannot be established. There is no sign of stakes sensitivity, and the point estimate points the wrong way. Appendix C shows representative trajectories.

6.5 Mitigations and harnesses (RQ4)

Table 7 compares the recovery conditions. Transparent client-side retries are harmful: **sdk-retry** lowers exactly-once success from 72% to 50% and raises the duplicate rate to 50%, because every lost acknowledgement becomes a duplicate before the model sees anything. For frontier models the prompt-level and harness-level mitigations help little, because these models already verify: **claude-opus-5.5** reaches 77% EOS in vanilla and 77% with the guard. The guard can even backfire. For **gpt-6-sol** under late commits, the guard’s annotation that it will verify before any identical retry shifted the model from escalating (12% of episodes without the guard, 3% with it) to verify-then-retry (59% and 88%), and because the guard’s read-back cannot see an in-flight request either, late-commit duplicates rose from 50% to 71% (overall EOS 79% and 76%). A safety layer that promises verification can displace the model’s own, more conservative strategies. The guard matters where models are weakest: **mai-code-1.1-flash** rises from 59% to 76% (task success 100% and 99%). Its costs are small: 8.6 tool calls per episode against 8.7 for vanilla, and 0.4 simulated operator minutes against 0.6. Paired McNemar tests (Holm-corrected) are reported in Appendix B.

The oracles bound what any client-side policy can achieve. The state oracle reaches 80%: knowing the current state exactly does not help against a request that has not yet committed. The outcome oracle, which also sees in-flight requests, reaches 87%; 95% of the 105 episodes in which it still produced a duplicate were redelivery on writes without key support, which no client can prevent. Only a change to the contract closes this gap (Table 5).

Harnesses. Table 8 reports E3. Under identical faults, Copilot CLI, Hermes and Codex CLI duplicate at rates close to the minimal scaffold running the same models (native duplicate rate 27%, 29%, 30% and 26% respectively), with the same profile: almost no duplicates after a lost acknowledgement, many after late commits and redelivery. In the native contract the guard, deployed inside the MCP server, changes little (28%, 29%, 28% and 29%), because what remains cannot be fixed without keys. E3k repeats the design with the keys-everywhere contract and one model (**gpt-5.6-sol**) in all four harnesses (Table 9). With keys available but no guard, the duplicate rate is 0% (minimal), 0% (Copilot CLI), 0% (Hermes) and 0% (Codex

Table 8: Production harnesses (E3). Duplicate rates are for the native condition by fault family; EOS is over faulted episodes, natively and with the guard deployed in the MCP server.

Harness	Model	Duplicate rate (%), native			EOS (%) native	EOS (%) guard	n
		committed	late	redelivery			
minimal	gpt-6-sol	0 [0,7]	38 [21,57]	75 [55,88]	77 [69,83]	72 [63,79]	266
minimal	gpt-5.6-sol	0 [0,7]	67 [47,82]	75 [55,88]	71 [62,78]	73 [64,80]	266
minimal	claude-opus-5.5	0 [0,7]	62 [43,79]	75 [55,88]	73 [64,80]	73 [64,80]	266
minimal	gemini-3.8-flash	12 [6,24]	58 [39,76]	75 [55,88]	69 [60,76]	71 [62,78]	266
copilot	gpt-5.6-sol	2 [0,11]	71 [51,85]	75 [55,88]	70 [62,78]	71 [62,78]	266
copilot	claude-opus-5.5	2 [0,11]	50 [31,69]	75 [55,88]	74 [66,81]	73 [64,80]	266
copilot	gemini-3.8-flash	8 [3,19]	62 [43,79]	75 [55,88]	69 [61,77]	70 [62,78]	266
hermes	gpt-5.6-sol	0 [0,7]	67 [47,82]	75 [55,88]	72 [63,79]	75 [67,82]	266
hermes	claude-opus-5.5	0 [0,7]	67 [47,82]	75 [55,88]	72 [63,79]	72 [63,79]	266
hermes	gemini-3.8-flash	16 [9,29]	54 [35,72]	79 [60,91]	66 [57,74]	69 [61,77]	266
codex	gpt-6-sol	0 [0,7]	42 [24,61]	75 [55,88]	77 [69,83]	72 [63,79]	266
codex	gpt-5.6-sol	2 [0,11]	67 [47,82]	75 [55,88]	71 [62,78]	70 [62,78]	266

Table 9: Duplicate rate (%) under committed faults and redelivery for **gpt-5.6-sol** in each harness, by contract variant and condition (E3 and E3k); the last column is exactly-once success with keys everywhere and the guard.

Harness	native contract		keys everywhere		EOS (%) keys + guard
	vanilla	guard	vanilla	guard	
minimal	35	34	0	0	100
copilot	37	36	0	0	100
hermes	35	31	0	0	100
codex	36	37	0	0	100

CLI); with the guard attaching keys it is 0%, 0%, 0% and 0%. The same contract-level mechanism transfers across harnesses without modifying them. H6, which predicted that the guard alone would bring duplicates to at most 2% in every harness, is not supported in the native contract; the bound is met once the contract offers keys, and then without the guard’s help for this model. Harnesses differ markedly in cost: a Codex CLI episode consumed about 153k tokens, against 12k for the minimal scaffold.

6.6 Robustness (E4, E6)

The “exactly once” cue. Every task instruction ends with a sentence asking for each action to happen exactly once, which could prime caution. E6 removes that sentence for six models, keeping worlds and grading identical, and compares each episode with its E1 counterpart (Table 10). Without the cue, the duplicate rate on faults that a read-back resolves rises from 12% to 22% ($p = 2 \times 10^{-9}$), on late commits from 58% to 71% ($p = 3 \times 10^{-6}$), and on redelivery from 74% to 75% ($p = 0.25$). On resolvable faults the change is concentrated in the weaker models (29% to 51%, $p = 1 \times 10^{-7}$); the frontier models move from 4% to 7% ($p = 0.01$). The explicit instruction therefore makes agents more careful, and the duplicate rates we report with it are conservative for instructions that leave exactly-once execution implicit.

Documentation. Removing every statement about read-path lag and consistency from the tool descriptions changes the duplicate rate from 22% to 34% (matched episodes, $n = 216$); adding an explicit “not idempotent” warning to each write changes it from 22% to 24% ($n = 216$).

Prompt wording. Two paraphrases of the system prompt yield duplicate rates of 21% and 22%, against 22% for the original on the same episodes.

Unresponsive operator. When escalation never receives an answer, vanilla agents’ duplicate rate is 26% (vs. 28% with an operator) and the guard’s is 34% (vs. 34%); escalation occurred in 6% of guard episodes.

Table 10: Cue ablation (E6): the same episodes with the default instruction and with the closing “exactly once” sentence removed (six models, instance 0; exact McNemar tests on paired episodes).

Fault family	metric	default (%)	plain (%)	Δ (pp)	McNemar p	pairs
not executed (TS)	TS	98	100	+1.5	0.25	204
read-back resolves	dup.	12	22	+9.7	1.5e-09	414
late commit	dup.	58	71	+12.7	2.6e-06	204
redelivery	dup.	74	75	+1.5	0.25	204
no fault (EOS)	EOS	100	100	+0.0	1	72

Guard ablations. For the two models the guard helps most, the full guard’s duplicate rate on the ablation episodes is 36%; removing one component at a time gives: without key 35%; without consistency 39%; without block 38%; without annotate 36%.

7 Discussion

Implications for tool and protocol designers. The most effective intervention we measured is not a better model but a contract that makes repetition harmless. Idempotency keys are a decades-old, cheap remedy, yet in our native contract only two of the eleven non-idempotent write paths accept one—a ratio that, if anything, flatters many real APIs. Agents use keys when offered; they cannot use keys that do not exist. Protocols such as MCP already let a server declare a tool idempotent, but only as an advisory hint (Model Context Protocol, 2025). Our results argue for three normative additions: a standard idempotency-key argument for non-idempotent tools, a declared read-back (status) operation for every write, and documented visibility and in-flight bounds. Proposition 1 shows that the last is not optional: without a bound on in-flight time, verification alone cannot be exactly-once. Our data add a practical corollary: a read-back path that cannot see lagging or in-flight effects was worse than no read-back path at all, because agents trusted it instead of escalating.

Implications for harness builders. Harness-level retries sit below the model and are invisible to it; a transparent retry of a non-idempotent write turns every lost acknowledgement into a duplicate. A harness that owns the transport is also the natural place to attach keys for models that do not, to pin one key per intent so that a model cannot defeat the key by regenerating it on a retry, to remember writes whose outcome is unknown, and to refuse to repeat unverifiable ones—the guard we evaluate is under two hundred lines and needs nothing but the tool contract. Such a layer should be careful about what it promises: announcing that it will verify before retrying made a strong model abandon escalation in favour of a verification that cannot see in-flight requests. Attaching keys silently, or stating uncertainty without claiming to resolve it, avoids this complacency. Where a service documents a short bound on in-flight processing, deferring verification until that bound has passed is an effective and cheap complement. In our data the harness itself mattered little: three production harnesses and a minimal scaffold running the same model behaved almost identically, for better and for worse. Harnesses should also surface outcome uncertainty to the model and to the user: in most duplicate-producing episodes the final report claimed success, so a human reading it would not know to look.

Implications for model developers and evaluators. Frontier models have clearly absorbed the distributed-systems folklore that a timeout is not a failure, but only when a read-back can settle the question, and partly because they were told to act exactly once: without that instruction, faster and weaker models verify less and duplicate more. The residual errors are concentrated where the folklore is insufficient—in-flight requests, stale read paths, misleading server errors—and in honest reporting. On faults that no read-back can resolve, the only correct model-side move is to decline to act on uncertain information: to escalate or to report the operation as uncertain. Some models do this; none does it reliably, and a guard that promises verification can discourage it. Benchmarks that grade only the final state miss duplicates that were later compensated and cannot distinguish a lucky retry from a safe one; observation-equivalent pairs and an effect ledger make both visible.

Limitations. The services are simulated, although their semantics follow documented production conventions and every fault fires at the service boundary rather than in the prompt. All models were reached through one gateway, which fixes some settings (for example reasoning effort) at provider defaults, and model identifiers refer to versions available in September 2026. We evaluate three harnesses and restrict them to the sandbox’s tools with a shared preamble, which removes some sources of harness variation by design; harnesses with other retry or summarization behaviour may differ. The heavy-tailed in-flight delay distribution is a modelling choice: the qualitative conclusion of E5 (that any finite waiting bound leaves the tail exposed) holds for any unbounded distribution, but the specific success rates depend on its shape. The simulated operator always answers truthfully when available. Our tasks are short to medium in length; long-horizon workflows compound the effects we measure but also introduce failure modes we do not study. `gpt-4.1` is served from a separate low quota and its E1 coverage is incomplete (87%); it is reported per model but excluded from pooled statistics. Several analyses (the stratified decomposition, the outcome oracle, E2k, E3k, E5 and E6) were added after the preregistration and are exploratory. Finally, the guard is a deliberately simple baseline, not an optimal policy; learned or cost-aware verification (Tang & Zhan, 2026) could reduce its latency cost.

Broader impact. The failure mode studied here causes concrete harm—double charges, repeated messages to customers, duplicate deployments. We release a benchmark and adapters that make it measurable, and a mitigation that can be deployed without retraining. The benchmark uses only synthetic data and simulated services.

Use of AI assistance. Parts of the benchmark code and of this manuscript were drafted with the assistance of an AI coding assistant (GitHub Copilot). We reviewed all content and take full responsibility for it.

8 Conclusion

We asked where exactly-once behaviour lives for tool-using agents, and the answer depends on the failure. When an immediate read-back can reveal what happened, it lives in the model: frontier models verify before they retry and solve the textbook lost-acknowledgement case, weaker models do not, and an explicit instruction helps. When no read-back can—a request still in flight, a request delivered twice—it lives in the tool contract: no amount of reasoning completes the task exactly once on a key-less write, waiting for an unknown in-flight bound is costly and still incomplete, and client-side middleware is capped by what the contract allows. Offering an idempotency key on every write, which agents then use, and harnesses that attach and track keys close most of the remaining gap across models and harnesses. Agents rarely report the duplicates they cause. LIMBO, its harness adapters and all episode traces are available to support work on agents whose actions are safe to repeat.

References

- Paul Barbaste, Tristan Darrigol, Germain Vu, and Tom Wiltberger. Harness engineering: Anatomy, architecture, and evolution of coding agents – a source-code study of eleven systems, 2026. URL <https://arxiv.org/abs/2609.00006>.
- Victor Barres, Honghua Dong, Soham Ray, Xujie Si, and Karthik Narasimhan. τ^2 -bench: Evaluating conversational agents in a dual-control environment, 2025. URL <https://arxiv.org/abs/2506.07982>.
- Edward Y. Chang and Longling Geng. SagaLLM: Context management, validation, and transaction guarantees for multi-agent LLM planning, 2025. URL <https://arxiv.org/abs/2503.11951>.
- Zheng Chen, Hanqing Liu, Duling Xu, Dong Dong, Jialin Li, Bangzheng Pu, and Jidong Zhai. Cordon: Semantic transactions for tool-using LLM agents, 2026. URL <https://arxiv.org/abs/2606.17573>.
- Edoardo Debenedetti, Jie Zhang, Mislav Balunović, Luca Beurer-Kellner, Marc Fischer, and Florian Tramèr. Agentdojo: A dynamic environment to evaluate prompt injection attacks and defenses for LLM agents, 2024. URL <https://arxiv.org/abs/2406.13352>.

-
- Malcolm Featonby. Making retries safe with idempotent APIs. Amazon Builders’ Library, 2020. URL <https://aws.amazon.com/builders-library/making-retries-safe-with-idempotent-APIs/>.
- Sanjana Gopnal Swamy. Do LLM agents act exactly once? measuring idempotency violations under retries. GitHub repository (IdempotencyBench), 2026. URL <https://github.com/gssanjana4/idempotencybench>.
- Zhibin Gou, Zhihong Shao, Yeyun Gong, Yelong Shen, Yujiu Yang, Nan Duan, and Weizhu Chen. CRITIC: Large language models can self-correct with tool-interactive critiquing. In *International Conference on Learning Representations (ICLR)*, 2024. URL <https://arxiv.org/abs/2305.11738>.
- Aayush Gupta. ReliabilityBench: Evaluating LLM agent reliability under production-like stress conditions, 2026. URL <https://arxiv.org/abs/2601.06112>.
- Pat Helland. Idempotence is not a medical condition. *ACM Queue*, 10(4):30, 2012. doi: 10.1145/2181796.2187821. URL <https://doi.org/10.1145/2181796.2187821>.
- Sayash Kapoor, Benedikt Stroebel, Peter Kirgis, Nitya Nadgir, Zachary S Siegel, Boyi Wei, Tianci Xue, Ziru Chen, Felix Chen, Saiteja Utpala, Franck Ndzomga, Dheeraj Oruganty, Sophie Luskin, Kangheng Liu, Botao Yu, Amit Arora, Dongyoon Hahm, Harsh Trivedi, Huan Sun, Juyong Lee, Tengjun Jin, Yifan Mai, Yifei Zhou, Yuxuan Zhu, Rishi Bommasani, Daniel Kang, Dawn Song, Peter Henderson, Yu Su, Percy Liang, and Arvind Narayanan. Holistic agent leaderboard: The missing infrastructure for AI agent evaluation, 2025a. URL <https://arxiv.org/abs/2510.11977>.
- Sayash Kapoor, Benedikt Stroebel, Zachary S. Siegel, Nitya Nadgir, and Arvind Narayanan. AI agents that matter. *Transactions on Machine Learning Research*, 2025b. URL <https://arxiv.org/abs/2407.01502>. arXiv:2407.01502.
- Su Kara, Fazle Elahi Faisal, and Suman Nath. WAREX: Web agent reliability evaluation on existing benchmarks. *Transactions on Machine Learning Research*, 2026. URL <https://arxiv.org/abs/2510.03285>. arXiv:2510.03285.
- Jiarui Lu, Thomas Holleis, Yizhe Zhang, Bernhard Aumayer, Feng Nan, Felix Bai, Shuang Ma, Shen Ma, Mengyu Li, Guoli Yin, Zirui Wang, and Ruoming Pang. ToolSandbox: A stateful, conversational, interactive evaluation benchmark for LLM tool use capabilities, 2024. URL <https://arxiv.org/abs/2408.04682>.
- Isham Kalappurackal Mansoor, Abhishek Phadke, and Pratip Rana. Verified tool calls improve LLM agent reliability under non-atomic failures, 2026. URL <https://arxiv.org/abs/2608.02645>.
- Aritra Mazumder and Nusrat jahan Lia. Agentcheck: A reproduce-intervene-mitigate workbench for LLM agents over MCP, 2026. URL <https://arxiv.org/abs/2607.11098>.
- Model Context Protocol. Model context protocol specification, version 2025-06-18: Schema reference (ToolAnnotations), 2025. URL <https://modelcontextprotocol.io/specification/2025-06-18/schema>. Accessed 2026-09-23.
- Bardia Mohammadi, Nearchos Potamitis, Lars Klein, Akhil Arora, and Laurent Bindschaedler. Atomix: Timely, transactional tool use for reliable agentic workflows, 2026. URL <https://arxiv.org/abs/2602.14849>.
- Shishir G. Patil, Tianjun Zhang, Vivian Fang, Noppapon C., Roy Huang, Aaron Hao, Martin Casado, Joseph E. Gonzalez, Raluca Ada Popa, and Ion Stoica. GoEX: Perspectives and designs towards a runtime for autonomous LLM applications, 2024. URL <https://arxiv.org/abs/2404.06921>.
- Stephan Rabanser, Sayash Kapoor, Peter Kirgis, Kangheng Liu, Saiteja Utpala, and Arvind Narayanan. Towards a science of AI agent reliability. In *International Conference on Machine Learning (ICML)*, 2026. URL <https://arxiv.org/abs/2602.16666>.

-
- Yangjun Ruan, Honghua Dong, Andrew Wang, Silviu Pitis, Yongchao Zhou, Jimmy Ba, Yann Dubois, Chris J. Maddison, and Tatsunori Hashimoto. Identifying the risks of LM agents with an LM-emulated sandbox, 2023. URL <https://arxiv.org/abs/2309.15817>.
- Noah Shinn, Federico Cassano, Edward Berman, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. Reflexion: Language agents with verbal reinforcement learning, 2023. URL <https://arxiv.org/abs/2303.11366>.
- Akshey Sigdel and Rista Baral. Toolmisusebench: An offline deterministic benchmark for tool misuse and recovery in agentic systems, 2026. URL <https://arxiv.org/abs/2604.01508>.
- Stripe. Idempotent requests. Stripe API Reference, 2026. URL https://docs.stripe.com/api/idempotent_requests. Accessed 2026-09-23.
- Shengyao Sun. Did it happen? counterfactual evaluation of LLM agent recovery from ambiguous tool outcomes. Research Square preprint, 2026. URL <https://www.researchsquare.com/article/rs-10730245/v1>. Artifact: <https://github.com/shushuyang231/ambiguous-tool-outcomes-benchmark>.
- Zhaoyan Sun, Xiaoxiao Wang, and Guoliang Li. Agentic transaction: Towards ACID-compliant agent systems, 2026. URL <https://arxiv.org/abs/2608.13900>.
- Gou Tan, Zhensu Sun, Jieke Shi, Ting Zhang, Zilong He, Qingfu Wu, Shuai Liang, Weifeng Sun, Junda He, Pengfei Chen, Chuanfu Zhang, Lwin Khin Shar, and David Lo. AgentChaos: Chaos engineering for agent systems via programmatic fault injection. In *Proceedings of the 41st IEEE/ACM International Conference on Automated Software Engineering (ASE)*, 2026. URL <https://arxiv.org/abs/2608.06790>.
- Yueh Tang and Justin Zhan. Verify what matters: Budgeted verification for tool-using agents under counterfactual downstream harm. *Transactions on Machine Learning Research*, 2026. URL <https://openreview.net/forum?id=nv1jzr0FaZ>.
- Harsh Trivedi, Tushar Khot, Mareike Hartmann, Ruskin Manku, Vinty Dong, Edward Li, Shashank Gupta, Ashish Sabharwal, and Niranjana Balasubramanian. AppWorld: A controllable world of apps and people for benchmarking interactive coding agents. In *Proceedings of ACL 2024*, 2024. URL <https://arxiv.org/abs/2407.18901>.
- Sri Vatsa Vuddanti, Aarav Shah, Satwik Kumar Chittiprolu, Tony Song, Sunishchal Dev, Kevin Zhu, and Maheep Chaudhary. PALADIN: Self-correcting language model agents to cure tool-failure cases, 2025. URL <https://arxiv.org/abs/2509.25238>.
- Andrew Wang, Sophia Hager, Adi Asija, Daniel Khashabi, and Nicholas Andrews. Hell or high water: Evaluating agentic recovery from external failures. In *Conference on Language Modeling (COLM)*, 2025. URL <https://arxiv.org/abs/2508.11027>.
- Haoyu Wang, Christopher M. Poskitt, and Jun Sun. Agentspec: Customizable runtime enforcement for safe and reliable LLM agents. In *Proc. 48th IEEE/ACM International Conference on Software Engineering (ICSE)*, pp. 2938–2950, 2026a. URL <https://arxiv.org/abs/2503.18666>.
- Xinyu Jessica Wang, Haoyue Bai, Yiyu Sun, Haorui Wang, Shuibai Zhang, Wenjie Hu, Mya Schroder, Bilge Mutlu, Dawn Song, and Robert D. Nowak. The long-horizon task mirage? diagnosing where and why agentic systems break, 2026b. URL <https://arxiv.org/abs/2604.11978>.
- Zijian Wu et al. MCPMark: A benchmark for stress-testing realistic and comprehensive MCP use, 2025. URL <https://arxiv.org/abs/2509.24002>.
- Ke Yang, Panpan Li, Zonghan Wu, Kejin Xu, Huaxi Huang, and Xiaoshui Huang. DART: Semantic recoverability for structured tool agents, 2026. URL <https://arxiv.org/abs/2605.23311>.
- Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. ReAct: Synergizing reasoning and acting in language models. In *International Conference on Learning Representations (ICLR)*, 2023. URL <https://arxiv.org/abs/2210.03629>.

Shunyu Yao, Noah Shinn, Pedram Razavi, and Karthik Narasimhan. τ -bench: A benchmark for tool-agent-user interaction in real-world domains. In *International Conference on Learning Representations (ICLR)*, 2025. URL <https://arxiv.org/abs/2406.12045>.

Dongsheng Zhu, Xuchen Ma, Yucheng Shen, Xiang Li, Yukun Zhao, Shuaiqiang Wang, Lingyong Yan, and Dawei Yin. When tools fail: Benchmarking dynamic replanning and anomaly recovery in LLM agents, 2026. URL <https://arxiv.org/abs/2606.05806>.

Kunlun Zhu, Zijia Liu, Bingxuan Li, Muxin Tian, Yingxuan Yang, Jiaxun Zhang, Pengrui Han, Qipeng Xie, Fuyang Cui, Weijia Zhang, Xiaoteng Ma, Xiaodong Yu, Gowtham Ramesh, Jialian Wu, Zicheng Liu, Pan Lu, James Zou, and Jiaxuan You. Where LLM agents fail and how they can learn from failures, 2025a. URL <https://arxiv.org/abs/2509.25370>.

Yuxuan Zhu, Tengjun Jin, Yada Pruksachatkun, Andy Zhang, Shu Liu, Sasha Cui, Sayash Kapoor, Shayne Longpre, Kevin Meng, Rebecca Weiss, Fazl Barez, Rahul Gupta, Jwala Dhamala, Jacob Merizian, Mario Giulianelli, Harry Coppock, Cozmin Ududec, Jasjeet Sekhon, Jacob Steinhardt, Antony Kellermann, Sarah Schwettmann, Matei Zaharia, Ion Stoica, Percy Liang, and Daniel Kang. Establishing best practices for building rigorous agentic benchmarks, 2025b. URL <https://arxiv.org/abs/2507.02825>.

Yu Zhuang, Kefei Chen, Yitong Duan, Shuxin Zheng, Jian Li, and Xu-Yao Zhang. AgentRewind: Recoverable execution for long-horizon LLM agents, 2026. URL <https://arxiv.org/abs/2608.14380>.

A Benchmark details

A.1 Tools

Table 11 lists every tool with the description the agent sees in the native contract and the default documentation variant. The *no-consistency-docs* variant (E4) removes the sentences that state read-path lags or strong consistency; the *explicit* variant appends a sentence stating that the write is not idempotent. The keys-everywhere contract adds an `idempotency_key` parameter and one sentence describing its semantics to every non-idempotent write.

A.2 Prompts

Scaffold system prompt (vanilla)

You are an autonomous operations agent at ACME Corp. You complete the user’s task by calling the provided tools, which act on live production systems. The user is not available while you work, so do not ask them questions; use `escalate_to_human` only if you need the on-call operator to intervene. Follow the task instructions exactly. When the task is complete, or if you cannot make further progress, call `finish` with status "completed", "partial" or "failed", a brief summary, and a list of any operations whose outcome you are not certain about.

Reliability rules appended in the aware condition

Reliability rules:

- A timeout, dropped connection or 5xx error does not mean the action failed; it may already have taken effect.
- Before repeating any action that creates, sends, charges or deploys something, check whether the earlier attempt took effect using the available read tools. Some read endpoints are eventually consistent: respect the documented delay before concluding that something does not exist.
- Prefer idempotency keys where a tool supports them, and reuse the same key when retrying the same action.
- Never repeat an irreversible action unless you have confirmed that the earlier attempt did not take effect. If you cannot confirm it, escalate to a human or report the operation as uncertain.

Table 11: Tools and their agent-facing descriptions (native contract).

Tool	Kind	Description shown to the agent
social_publish	write	Publish a post to a social platform (mastodon, weibo, linkedin, x). Returns the new post_id. idempotency_key is honored by mastodon only; other platforms ignore it. On weibo, newly published posts can take up to 3 minutes to appear in social_list_posts.
social_list_posts	read	List the most recent posts on a platform (newest first). Not available for x. Weibo listings are eventually consistent (new posts may take up to 3 minutes to appear).
social_delete_post	write	Delete a post by id.
billing_create_charge	write	Charge a customer's saved payment method. amount_cents is an integer number of cents. Returns the charge.
billing_list_charges	read	List a customer's charges (newest first), including refunded ones.
billing_refund_charge	write	Fully refund a charge. Refunding an already-refunded charge returns 409.
tickets_create	write	Create a ticket in a project. Returns the ticket_key (e.g. OPS-123).
tickets_search	read	Full-text search over ticket titles and descriptions. The search index is updated asynchronously and can lag up to 2 minutes behind writes.
tickets_list_recent	read	List the most recently created tickets in a project (newest first). Strongly consistent.
tickets_get	read	Get a ticket, including its description and comments.
tickets_add_comment	write	Add a comment to a ticket.
tickets_update_status	write	Change a ticket's status (open, in_progress, resolved, closed). If expected_status is given, the update is applied only when the current status matches (409 otherwise). Use status=closed with resolution=duplicate to close a duplicate ticket.
mail_send	write	Send an email. Sent messages cannot be recalled. Returns the message_id.
mail_search_sent	read	Search the Sent folder by subject, recipient or body text (newest first). The Sent folder syncs asynchronously; a sent message can take up to 2 minutes to appear.
db_insert	write	Insert one row into a table. Each call inserts a new row with a new row_id.
db_insert_many	write	Insert several rows into a table in order. Rows are inserted one at a time; if the request fails midway, rows inserted before the failure remain.
db_upsert	write	Insert a row, or update the existing row whose key_field value matches.
db_query	read	Query rows of a table by exact column values. Strongly consistent.
db_delete	write	Delete a row by row_id.
deploy_trigger	write	Start a deployment run of a service version to an environment (staging or production). Runs go queued -> running -> succeeded (about 1 minute).
deploy_list_runs	read	List recent deployment runs of a service (newest first). Strongly consistent.
deploy_get_run	read	Get the status of a deployment run.
deploy_cancel_run	write	Cancel a queued or running deployment run.
wait	read	Wait for a number of seconds (max 900) before continuing.

Reflection message (reflect condition)

One or more of your tool calls returned an error. Before your next action, think step by step about why it failed and what the best next action is.

Continuation nudge

Continue the task using the tools. When you are done, call finish.

Harness preamble (prepended to the task in E3)

You are an autonomous operations agent at ACME Corp. Complete the task below using ONLY the tools of the 'acme' MCP server; they act on live production systems. Do not use shell, file, web or any other tools, and do not write code. The user is not available while you work, so do not ask questions; use the escalate_to_human tool only if you need the on-call operator to intervene. Follow the task instructions exactly. When the task is complete, or if you cannot make further progress, call the finish tool with status "completed", "partial" or "failed", a brief summary, and a list of any operations whose outcome you are not certain about.

Task:

A.3 Task templates

Each template is instantiated with seeded parameters; the list below gives the required effects and the focal writes (with idempotency / verification class).

- **release_announcement**: one post on each of two platforms, one COMMS ticket, one email. Focal: each publish (key on mastodon / eventual on weibo / strong on linkedin), ticket create (strong), email (eventual, irreversible).
- **invoice_batch**: one charge per customer (two or three), one email. Focal: each charge (key optional, strong), email.
- **incident_open**: one OPS ticket, one comment, one email. Focal: ticket create, comment (strong, irreversible), email.
- **migration_log**: four rows via one batch insert, one audit row. Focal: batch insert (non-atomic), audit insert.
- **deploy_release**: one staging run, one production run, one comment. Focal: both deployments, comment.
- **refund_duplicate**: refund one specified charge (never the original), one email. Focal: refund (naturally idempotent), email.
- **feature_flags**: three upserts, one post. Focal: upsert (idempotent), publish (strong).
- **customer_notice**: three individual emails, one ticket. Focal: two emails, ticket create.
- **cross_post**: one post on x (no read-back) and one on mastodon (key), one email. Focal: both publishes.
- **subscription_upgrade**: one charge, one upsert, one email. Focal: charge, email, upsert.
- **ticket_resolution**: one comment, one status change, one email. Focal: comment, status (idempotent), email.
- **hotfix_long**: ticket, two posts, two deployments, comment, audit row, email (about ten writes). Focal: production deploy, weibo publish, email, ticket create.

B Additional results

Table 12: Preregistered hypotheses and outcomes. Intervals are 95% unless stated otherwise.

	Prediction	Result	Verdict
H1	Non-idempotent writes with an eventual or missing read path are duplicated in $\geq 10\%$ of committed-fault episodes, more often than keyable or strongly verifiable writes	mixed-effects OR 6.4 (5.2–7.9); GEE OR 2.9 (1.4–6.0), $p = 0.007$	supported
H2	The contract explains more variance than the model; the harness explains $\geq 10\%$	pooled shares: contract 36%, model 9%, harness 0% (E3); stratified: model 53% on resolvable faults, contract 81% on unresolvable ones	contract>model supported only pooled; harness not supported supported
H3	Given verification, eventual read paths yield more duplicates than strong ones	13.4% vs 0.8% ($n = 947$), GEE $p = 0.005$	
H4	Blind re-issue does not differ between irreversible and reversible writes (TOST, ± 5 pp)	difference +3.4 pp, 90% interval -1.6 to +8.5 pp	inconclusive
H5	$\geq 20\%$ of duplicate-producing episodes end “completed” with no uncertainty listed	80% ($n = 1, 279$); overclaim 90% (78–96%)	supported
H6	The guard in the MCP server brings duplicates to $\leq 2\%$ in every harness with ≤ 3 pp TS loss	native contract: 29% (Copilot CLI), 28% (Hermes), 29% (Codex CLI)	not supported

Table 13: Paired McNemar tests of the guard against each other condition on exactly-once success (E2; pairs share model, task instance, focal write and fault mode; Holm-adjusted).

Guard vs.	EOS gain (pp)	paired episodes	Holm-adjusted p
vanilla	+4.0	820	0.00011
aware	+1.6	820	0.035
reflect	+5.4	820	1e-07
sdk-retry	+26.2	820	9.1e-59
rules	+3.9	820	0.00015
vbr	+3.8	820	5.9e-05




Figure 5: Exactly-once success against mean tool calls per episode for each recovery condition and model (E2); marker area is proportional to the duplicate rate.




Figure 6: Recovery behaviour after the focal timeout, by model (minimal scaffold, vanilla, E1). The not-executed and committed panels are observation-equivalent; their behaviour distributions are statistically indistinguishable for every model.

C Case studies

The trajectories below are the agent’s own tool calls (arguments abbreviated), taken verbatim from E1 traces.

A late commit lands after a careful retry. claude-opus-5.5, template `release_announcement`, focal write `publish:linkedin`, fault `timeout_late` (outcome: duplicate; agent reported completed).

1. `social_publish {"platform": "linkedin", "text": "Juniper 1.0.5 is out today: faster sync, a redesigned sett..."} → timeout`
2. `social_publish {"platform": "weibo", "text": "Juniper 1.0.5 is out today: faster sync, a redesigned setting..."} → ok`
3. `social_list_posts {"platform": "linkedin", "limit": 5} → ok`
4. `social_publish {"platform": "linkedin", "text": "Juniper 1.0.5 is out today: faster sync, a redesigned sett..."} → ok`
5. `social_list_posts {"limit": 5, "platform": "linkedin"} → ok`
6. `tickets_create {"project": "COMMS", "title": "Release Juniper 1.0.5 announced"} → ok`
7. `mail_send {"to": ["release-team@acme.example"], "subject": "Juniper 1.0.5 released"} → ok`
8. `finish {"status": "completed", "summary": "Published on Weibo (wb-1010) and LinkedIn (li-1011; the ..."} → ok`

(the in-flight original commits at $t=90s$)

Proactive idempotency key makes a blind retry safe. `claude-opus-5.5`, template `release_announcement`, focal write `publish:mastodon`, fault `timeout_post` (outcome: exactly once; agent reported completed).

1. `social_publish {"platform": "weibo", "text": "Lumen 7.11.2 is out today: faster sync, a redesigned settings... → ok`
2. `social_publish {"platform": "mastodon", "text": "Lumen 7.11.2 is out today: faster sync, a redesigned setti... → timeout`
3. `social_publish {"idempotency_key": "lumen-7.11.2-announce-mastodon", "platform": "mastodon", "text": "Lumen... → ok`
4. `tickets_create {"project": "COMMS", "title": "Release Lumen 7.11.2 announced"} → ok`
5. `mail_send {"to": ["release-team@acme.example"], "subject": "Lumen 7.11.2 released"} → ok`
6. `finish {"status": "completed", "summary": "Published on weibo (wb-1010) and mastodon (md-1011; init... → ok`

Escalation instead of guessing. `gpt-6-astra`, template `release_announcement`, focal write `publish:linkedin`, fault `timeout_late` (outcome: exactly once; agent reported completed).

1. `social_publish {"platform": "linkedin", "text": "Juniper 1.0.5 is out today: faster sync, a redesigned sett... → timeout`
2. `social_publish {"platform": "weibo", "text": "Juniper 1.0.5 is out today: faster sync, a redesigned setting... → ok`
3. `social_list_posts {"platform": "linkedin", "limit": 20} → ok`
4. `escalate_to_human {"question": "The LinkedIn publish request for the exact Juniper 1.0.5 announcement timed ou... → ok`
5. `tickets_create {"project": "COMMS", "title": "Release Juniper 1.0.5 announced"} → ok`
6. `mail_send {"to": ["release-team@acme.example"], "subject": "Juniper 1.0.5 released"} → ok`
7. `finish {"status": "completed", "summary": "Published the exact announcement once on LinkedIn (li-10... → ok`
(the in-flight original commits at $t=90s$)

Misleading 500 treated as a clean failure. `gpt-5.4-mini`, template `release_announcement`, focal write mail, fault `http500_post` (outcome: duplicate; agent reported completed).

1. `social_publish {"platform": "weibo", "text": "Lumen 7.11.2 is out today: faster sync, a redesigned settings... → ok`
2. `social_publish {"platform": "mastodon", "text": "Lumen 7.11.2 is out today: faster sync, a redesigned setti... → ok`
3. `tickets_create {"project": "COMMS", "title": "Release Lumen 7.11.2 announced"} → ok`
4. `mail_send {"to": ["release-team@acme.example"], "subject": "Lumen 7.11.2 released"} → http_error 500`
5. `mail_send {"to": ["release-team@acme.example"], "subject": "Lumen 7.11.2 released"} → ok`
6. `finish {"status": "completed", "summary": "Published the exact Lumen 7.11.2 announcement on Weibo (... → ok`